

# Transformers

Machine Learning Course - CS-433

Nov 12, 2024

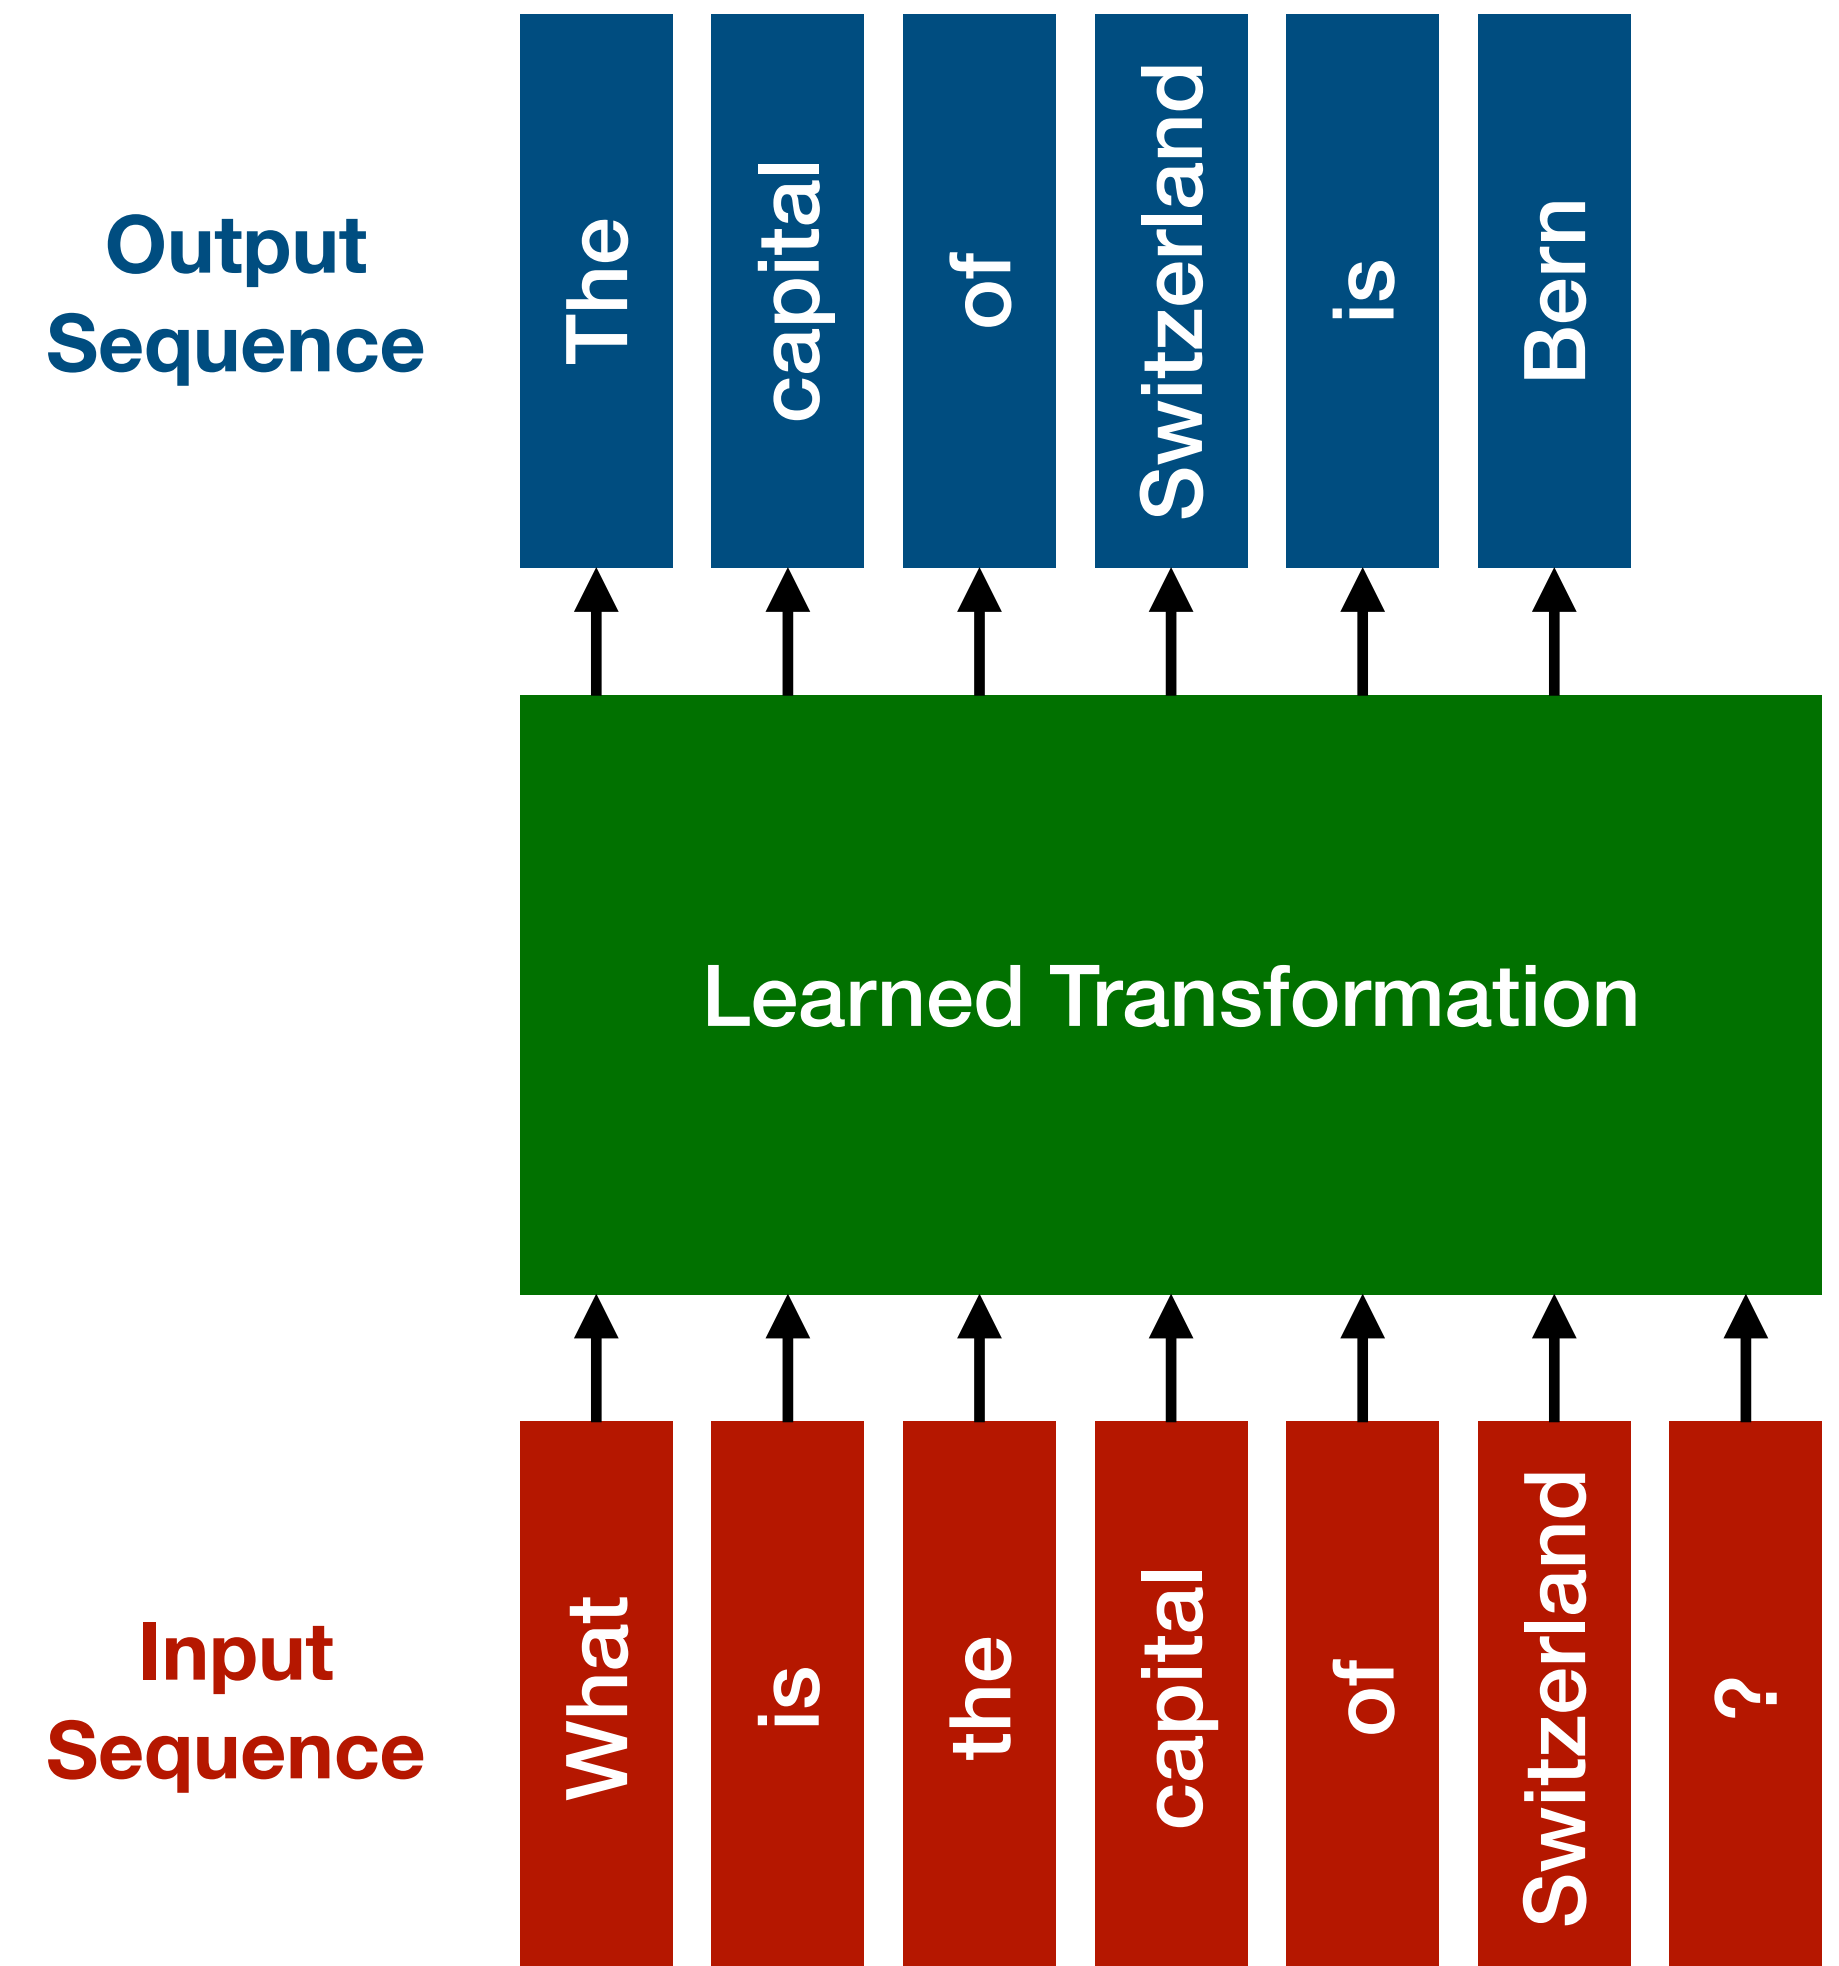
Nicolas Flammarion

**EPFL**

# Sequence-to-Sequence Transformations

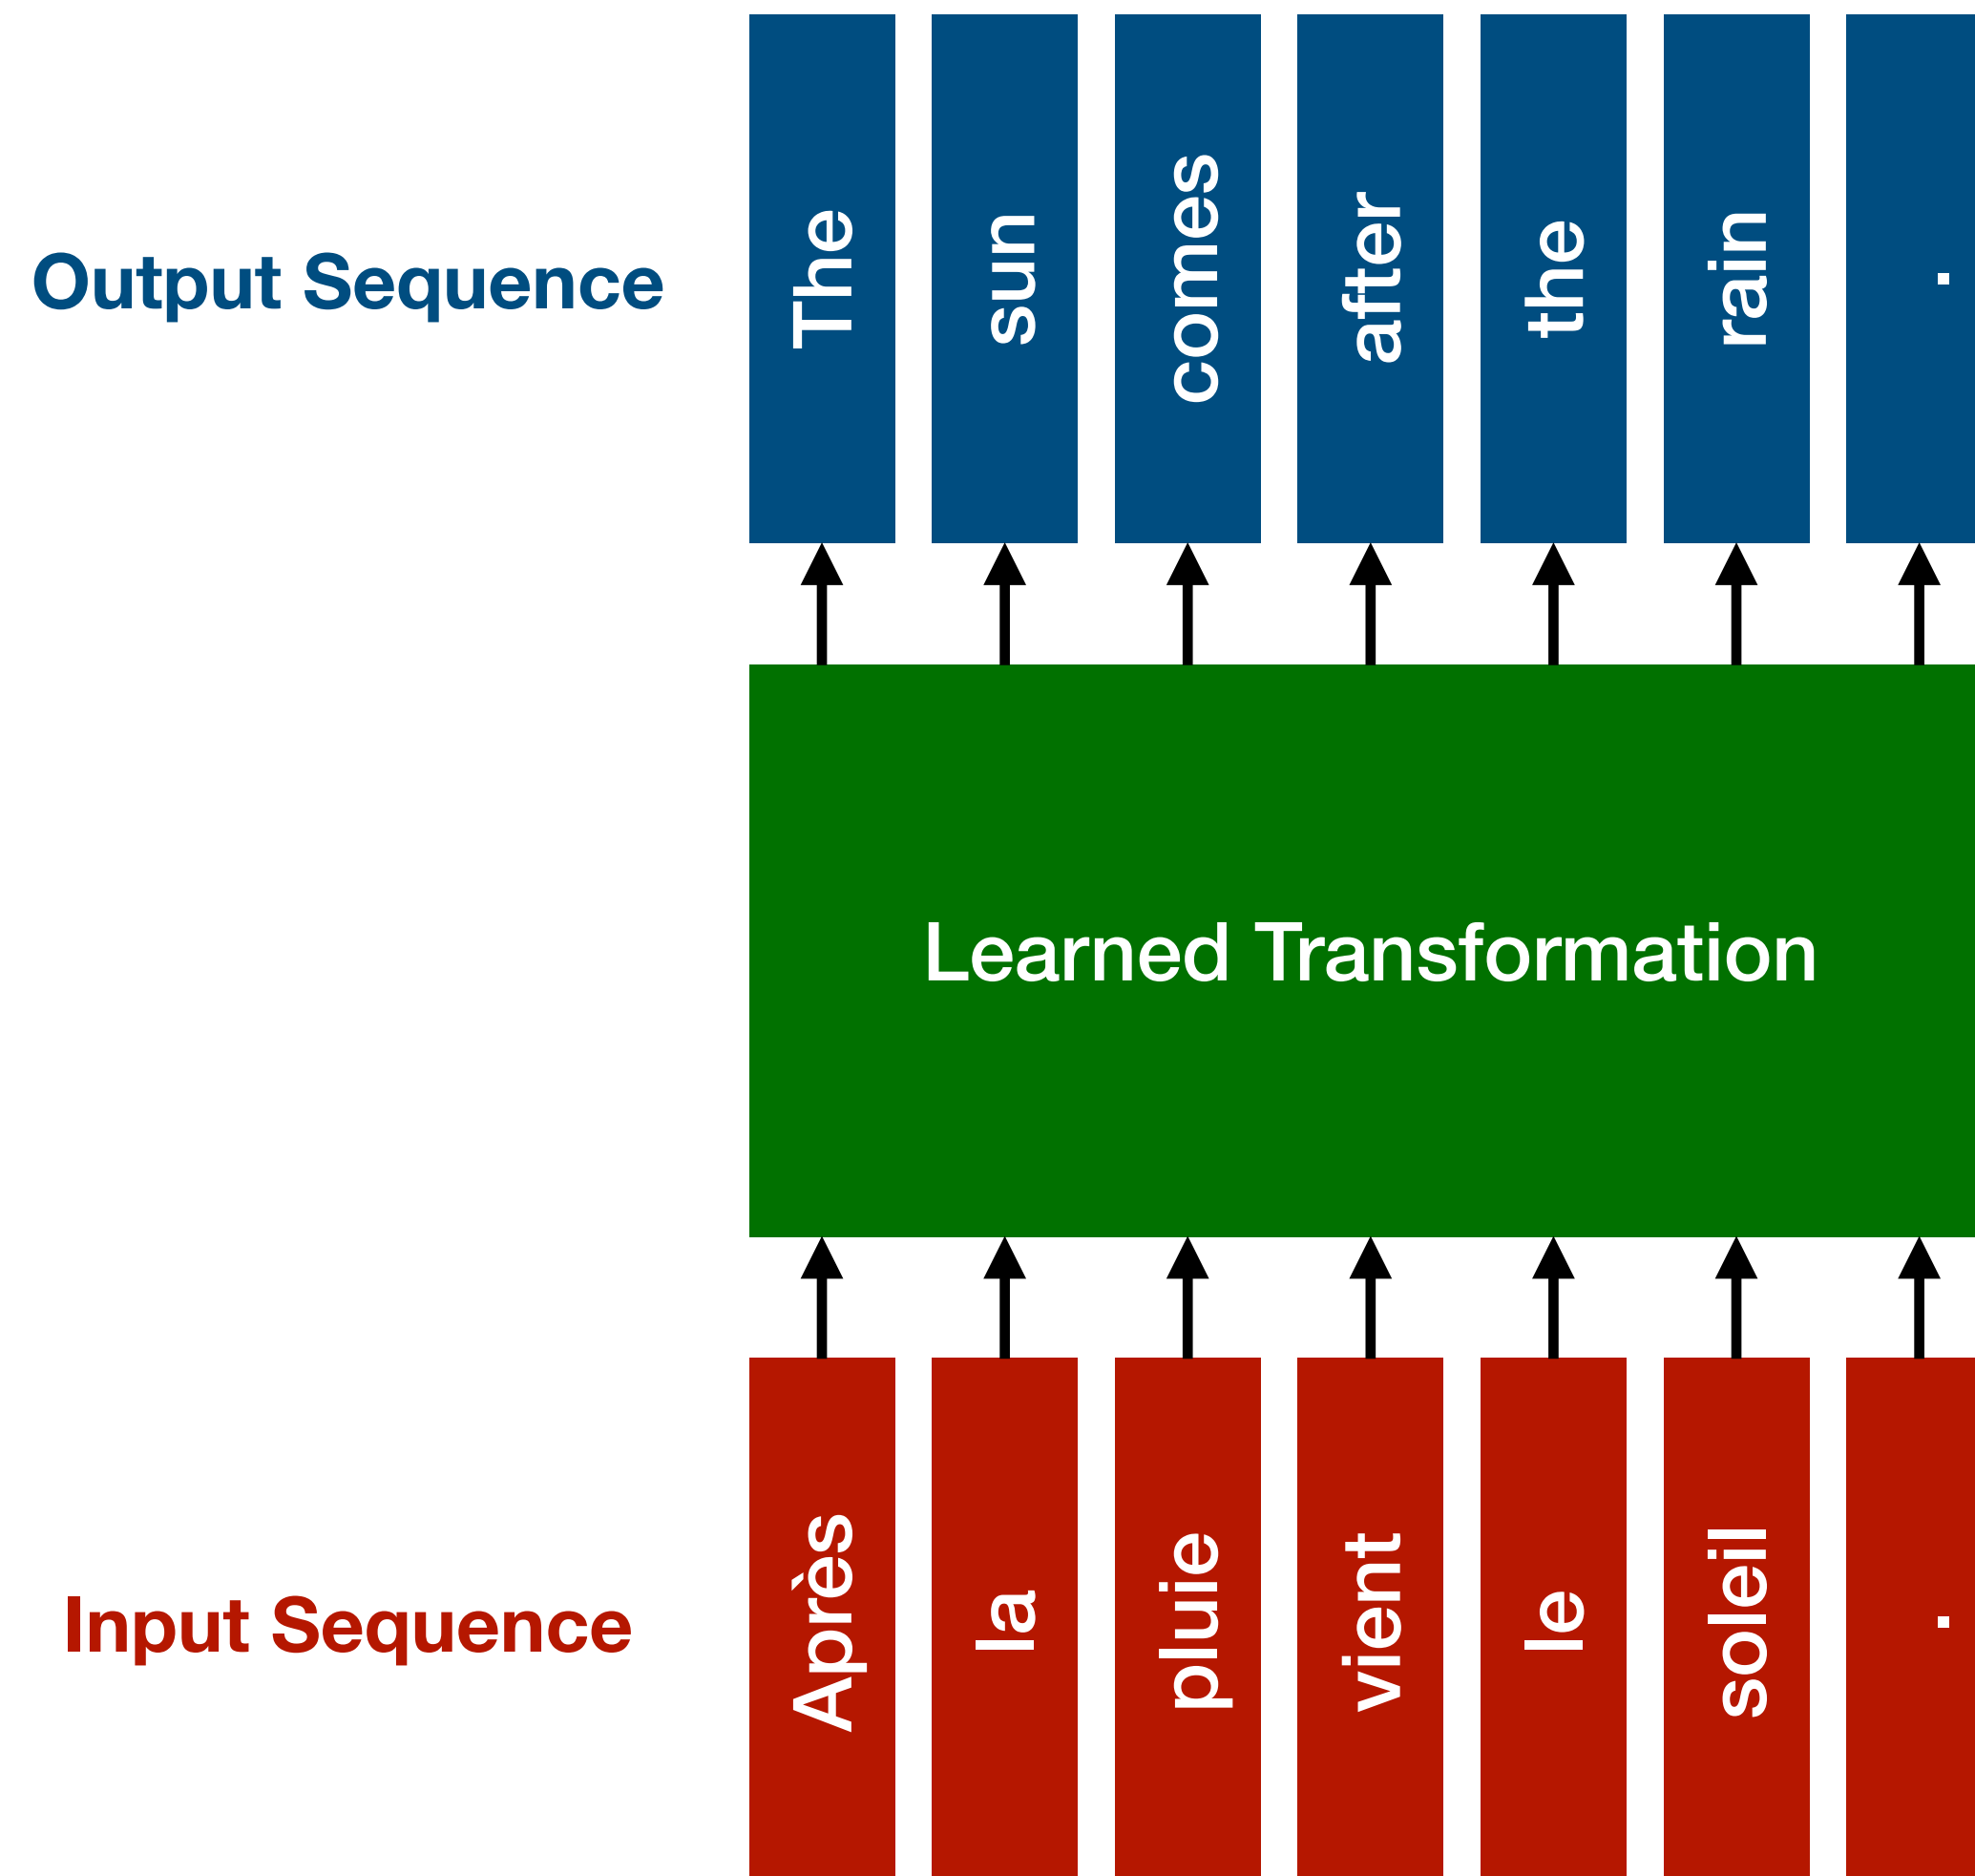
# Sequence-to-Sequence Transformations

- Many interesting problems in ML can be expressed as **mapping one sequence to another**
- **Example:** chatbots like ChatGPT
  - Input: question (word sequence)
  - Output: answer (word sequence)
- Input and output sequences can represent various types of data: words, images, speech, proteins, time series, etc.

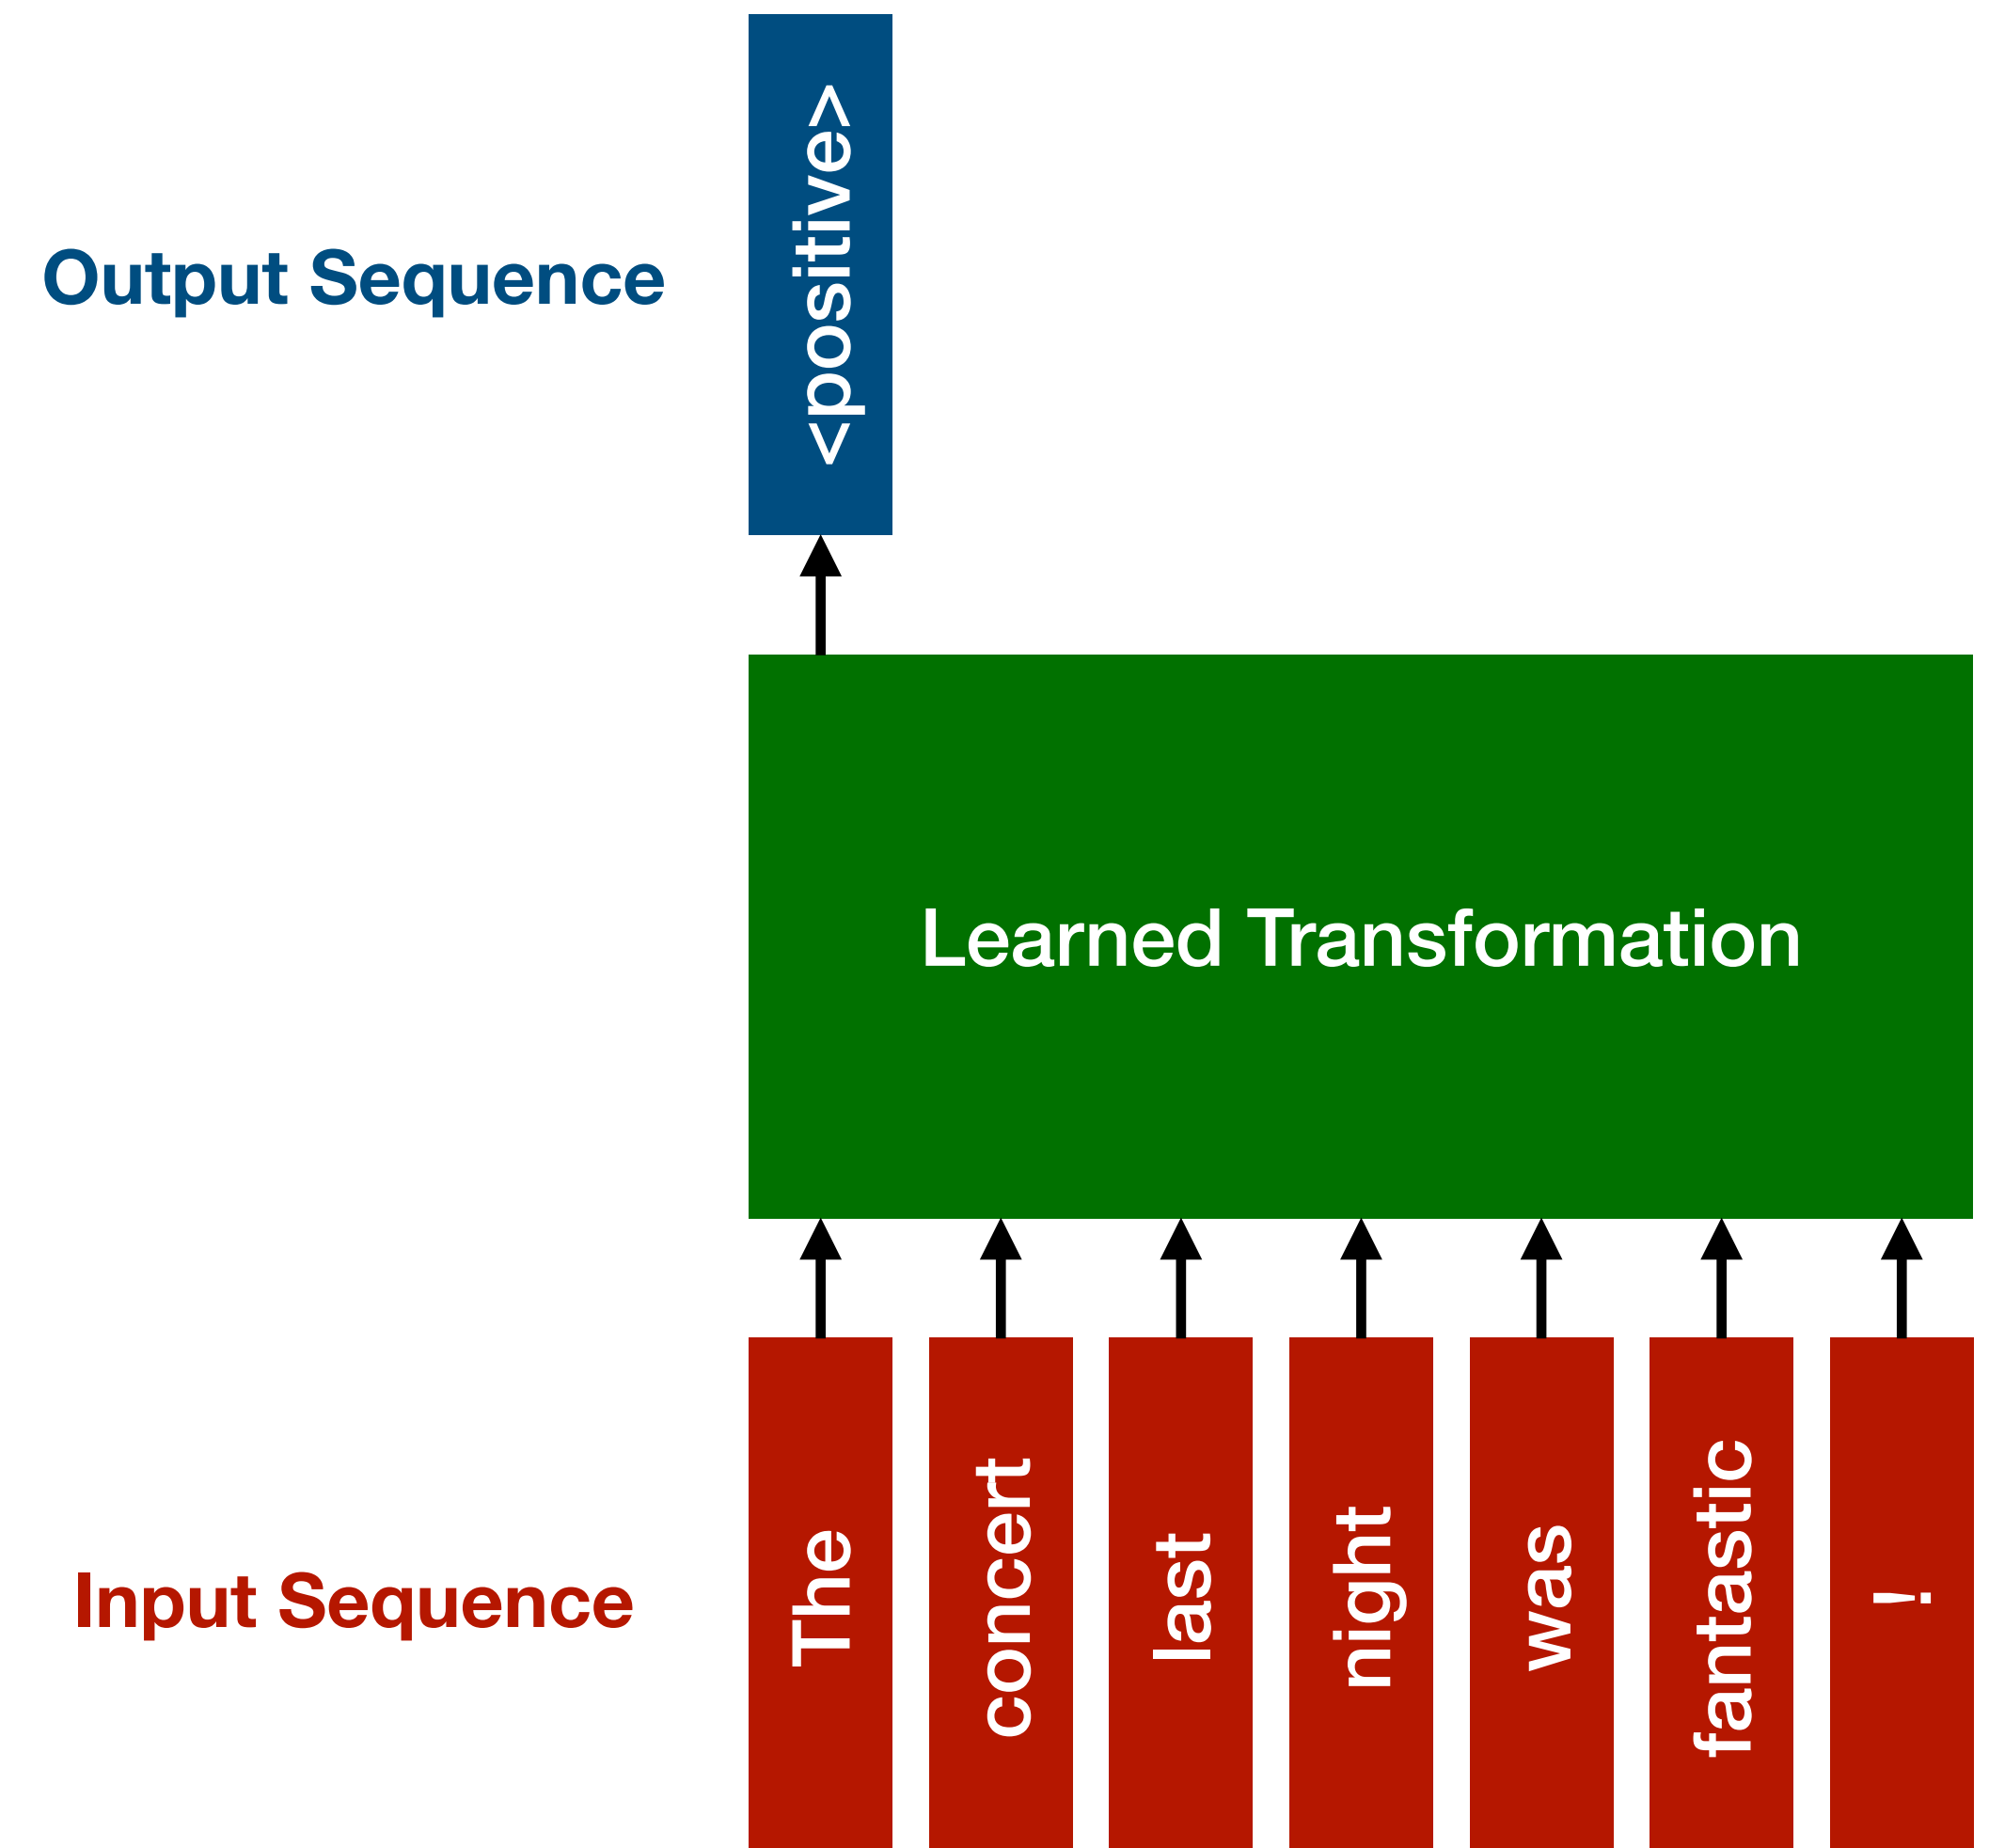


# The sequence-to-sequence framework is very general

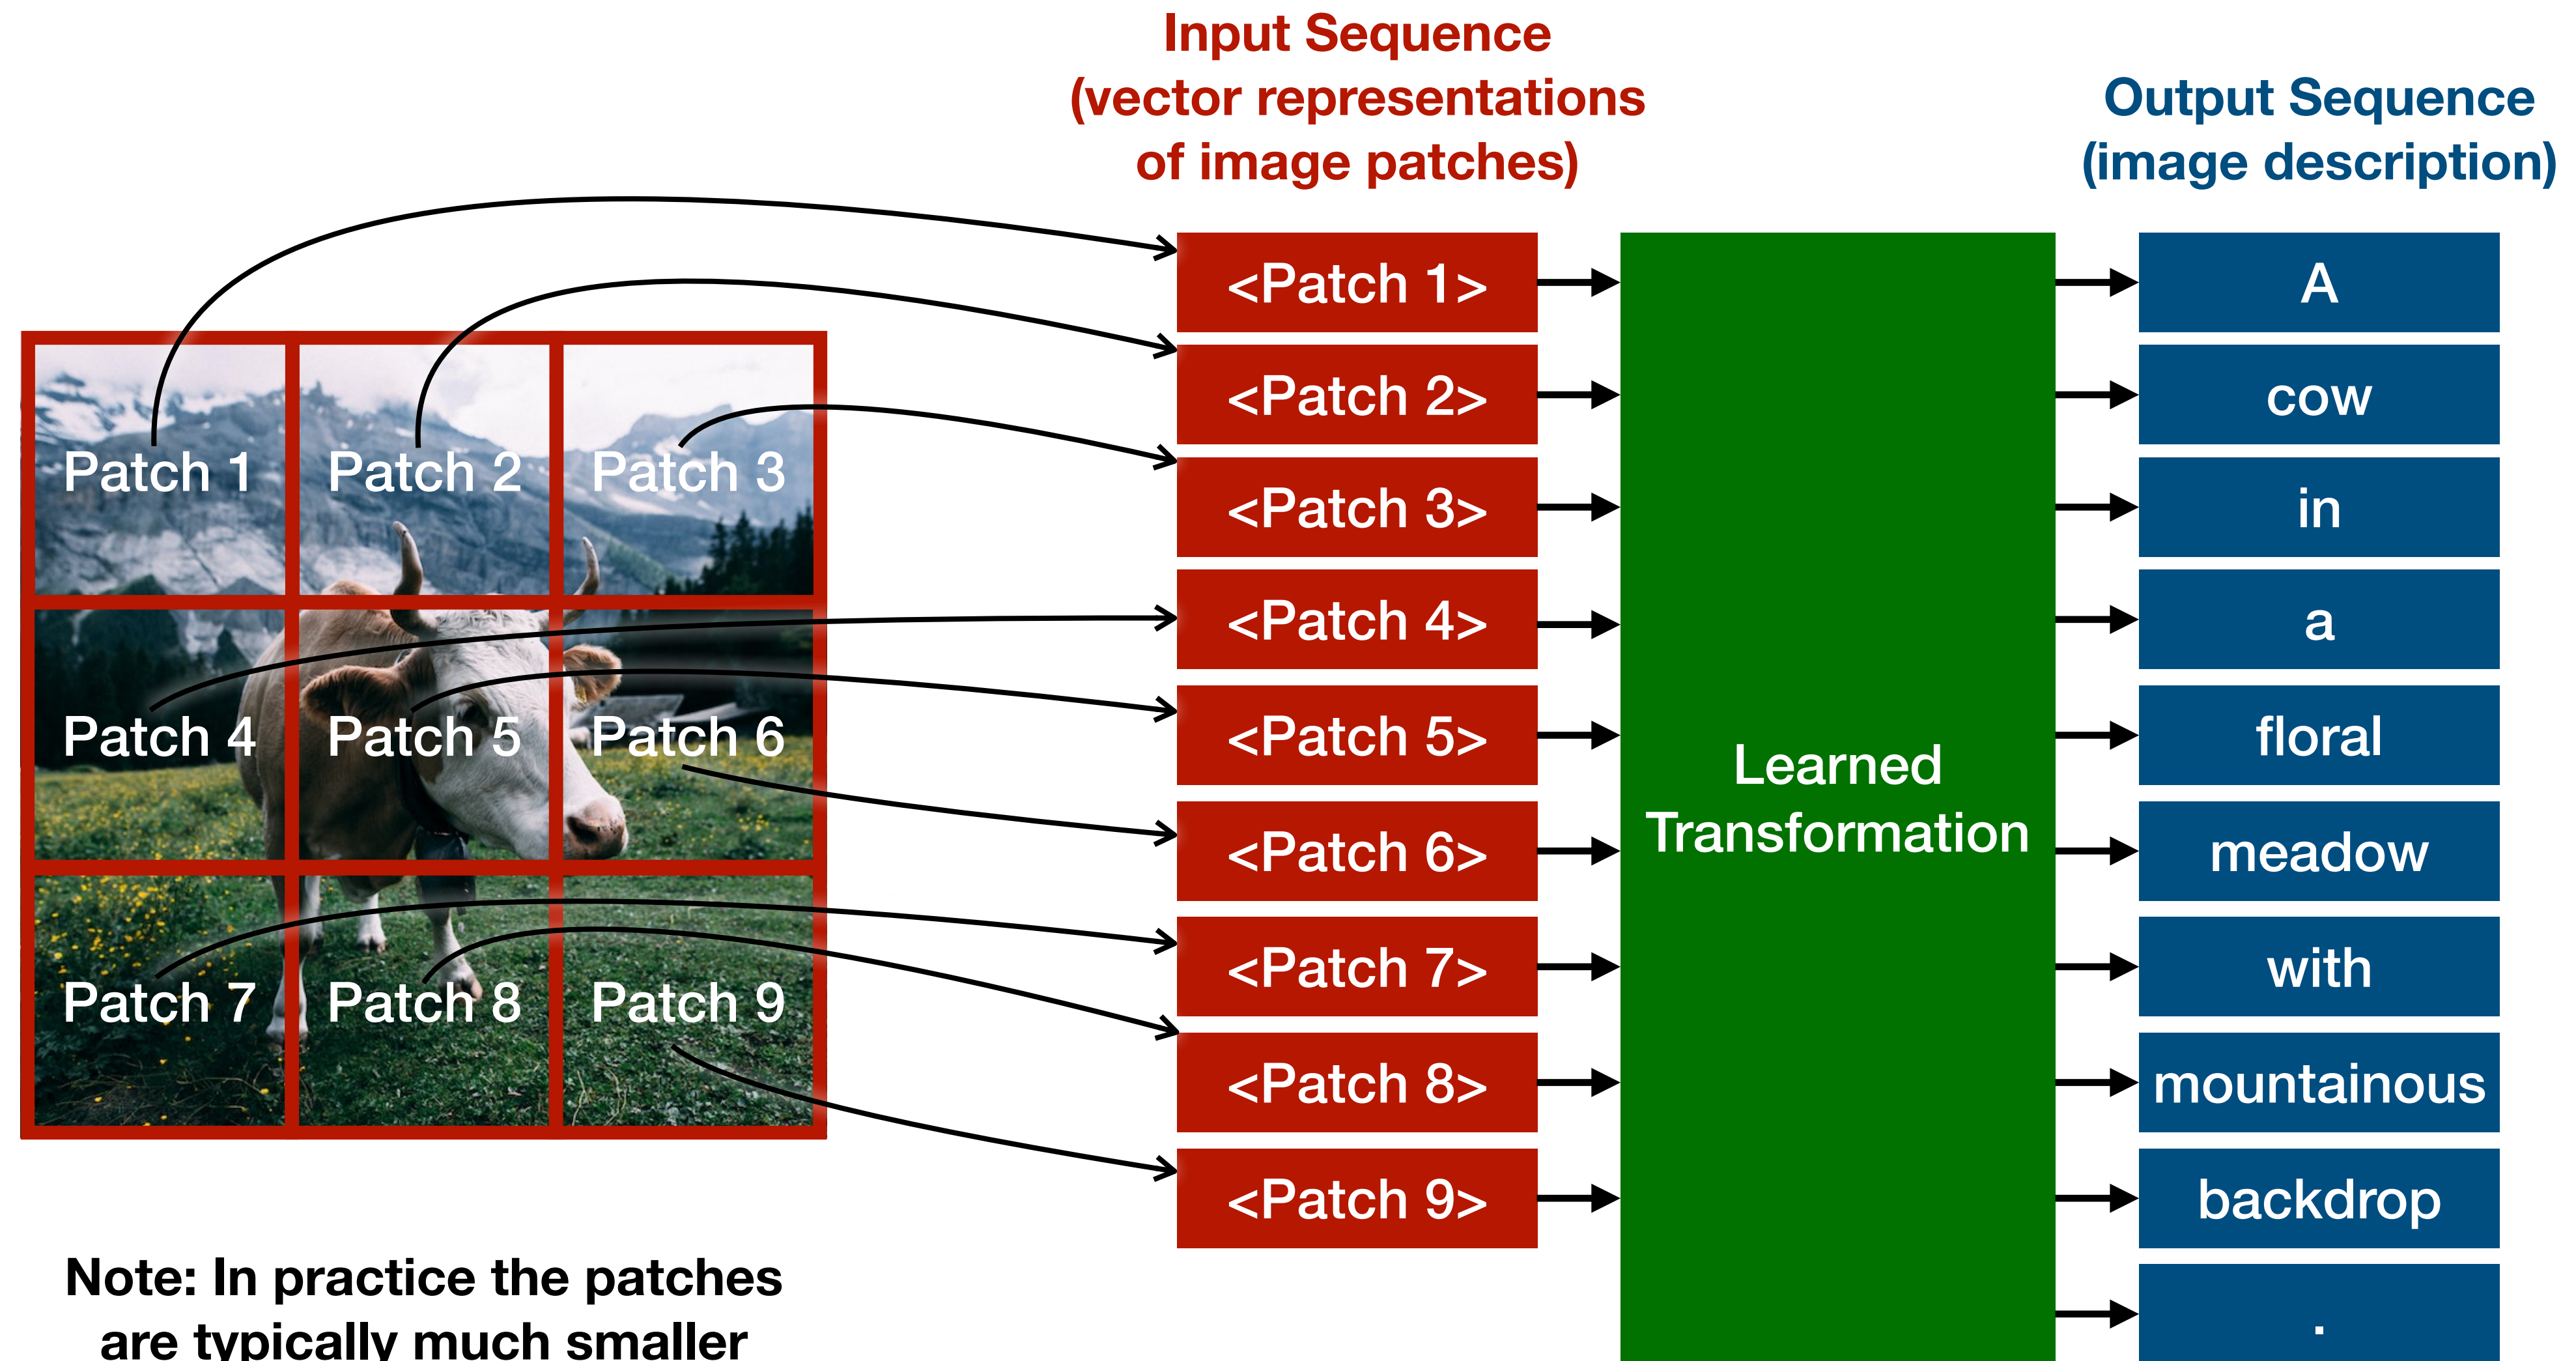
**Translation**



**Sentiment classification**  
(output is a sequence of length 1)

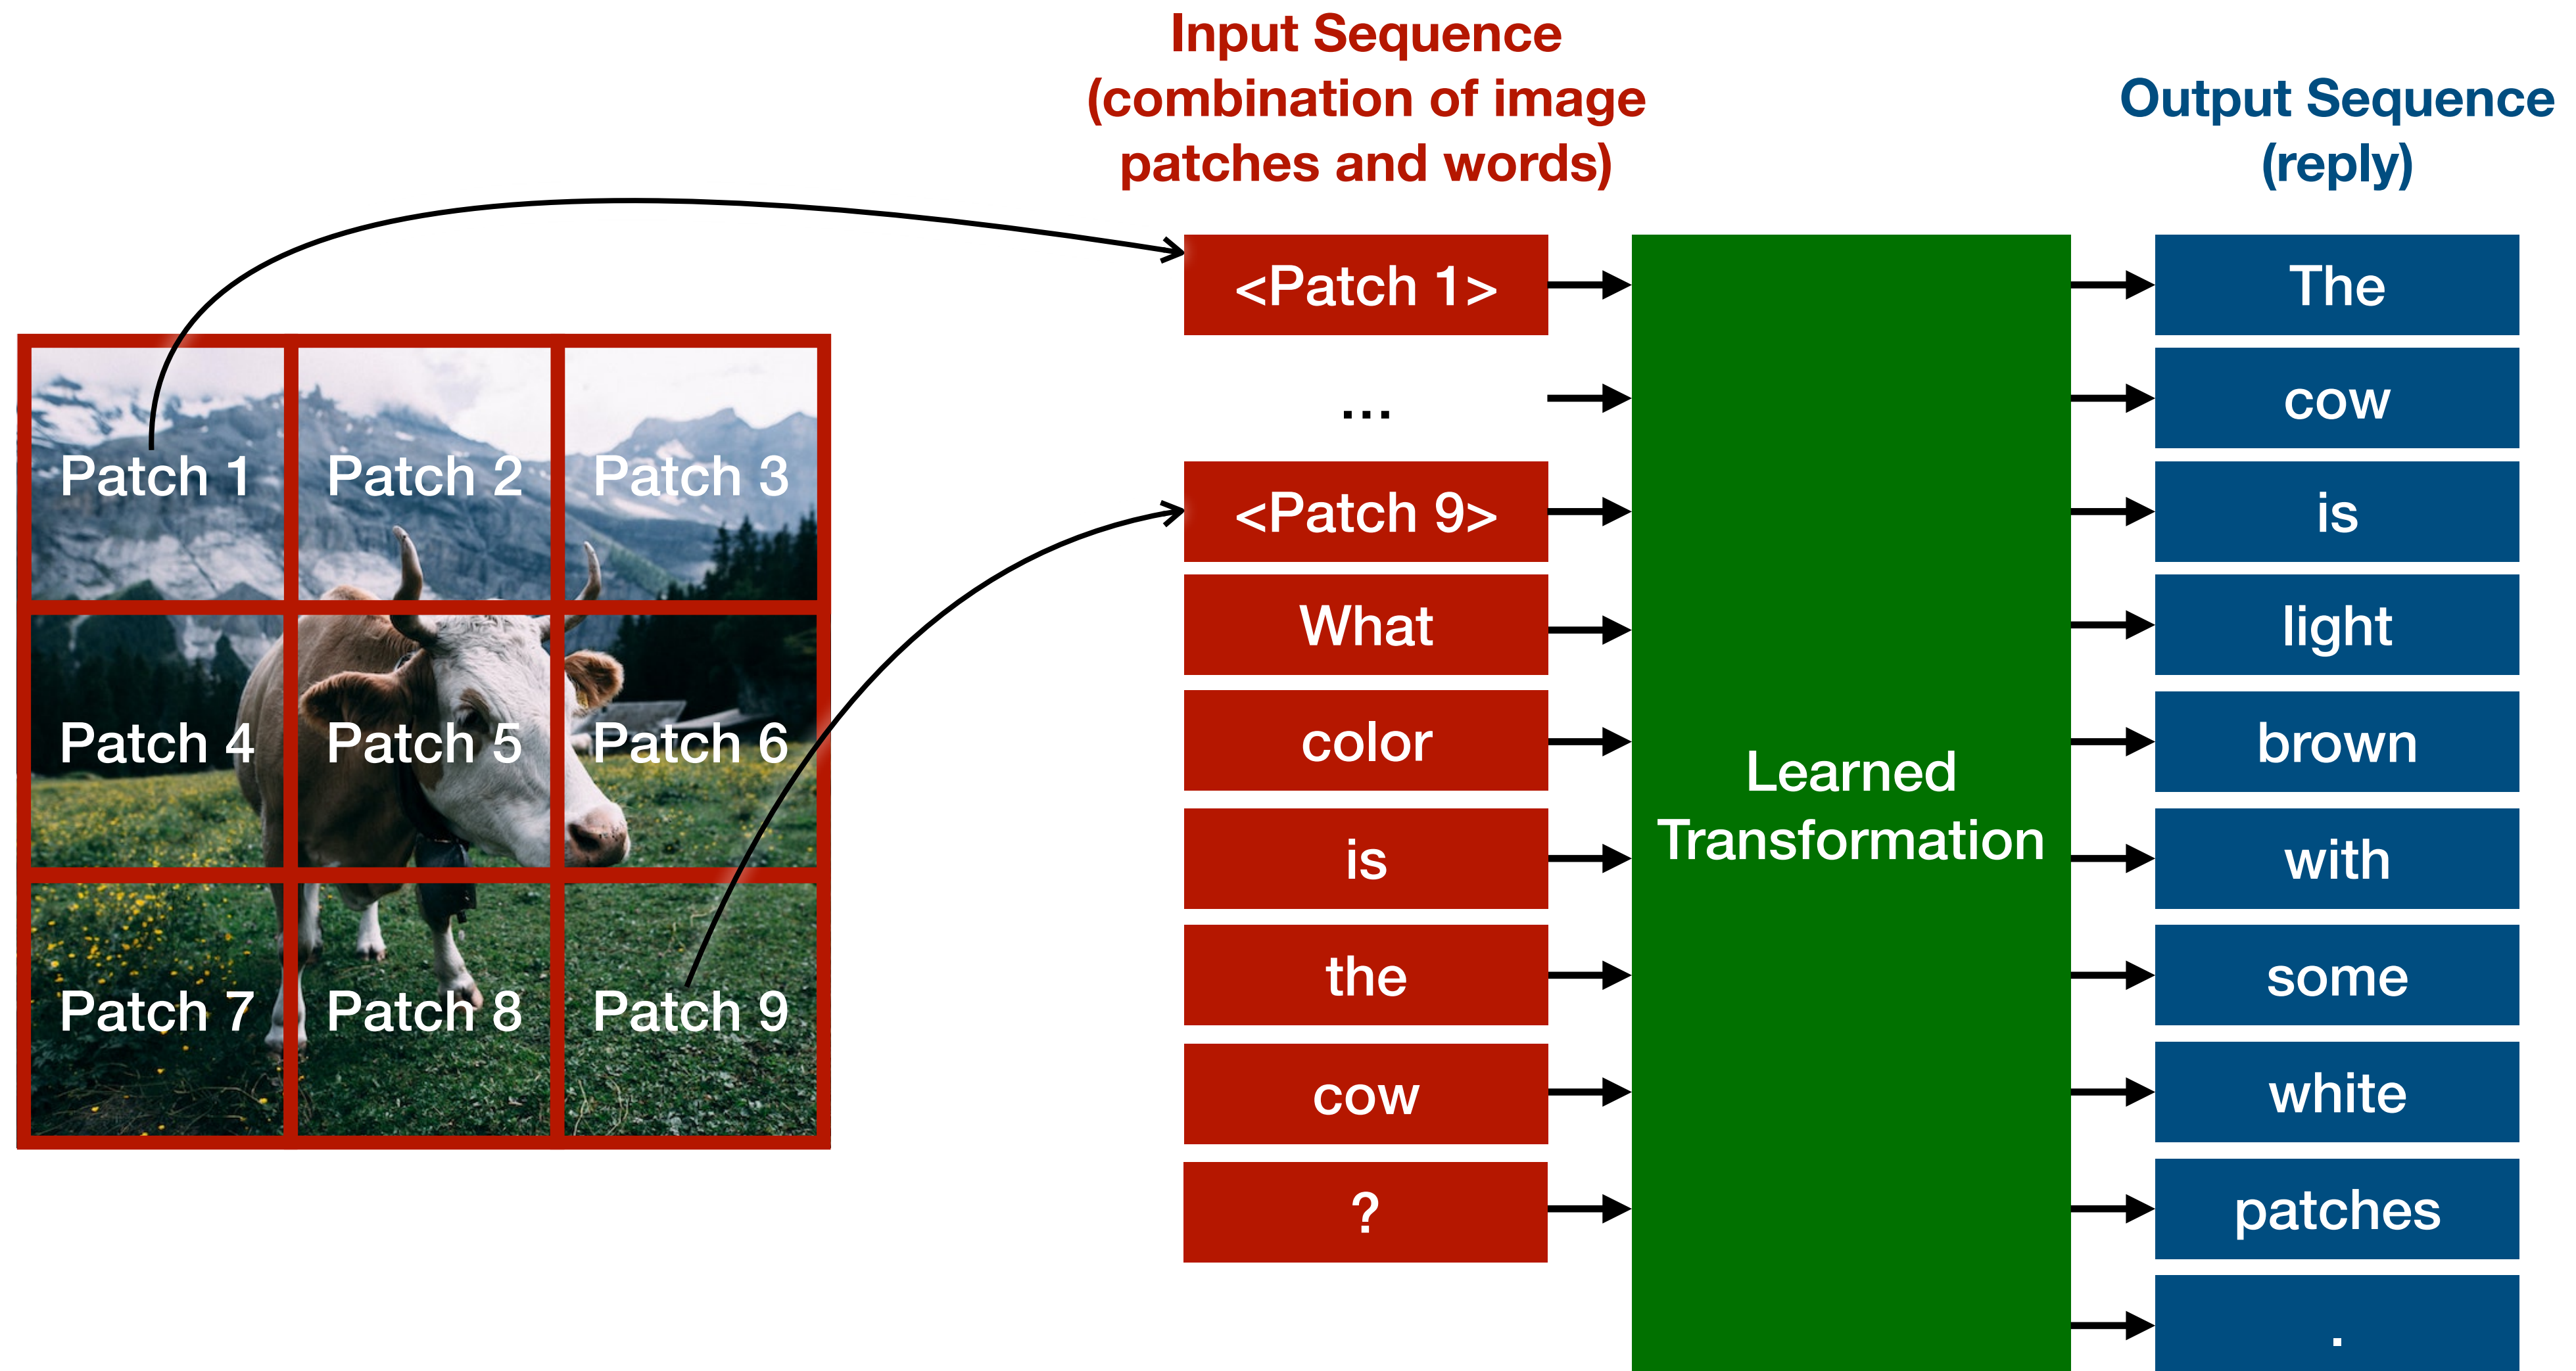


# Images can also be represented as sequences





# Sequences can be multimodal (image + text)



# Transformers

# What is a Transformer?

$$f : \textit{sequence} \rightarrow \textit{sequence}$$

(using **self-attention**)

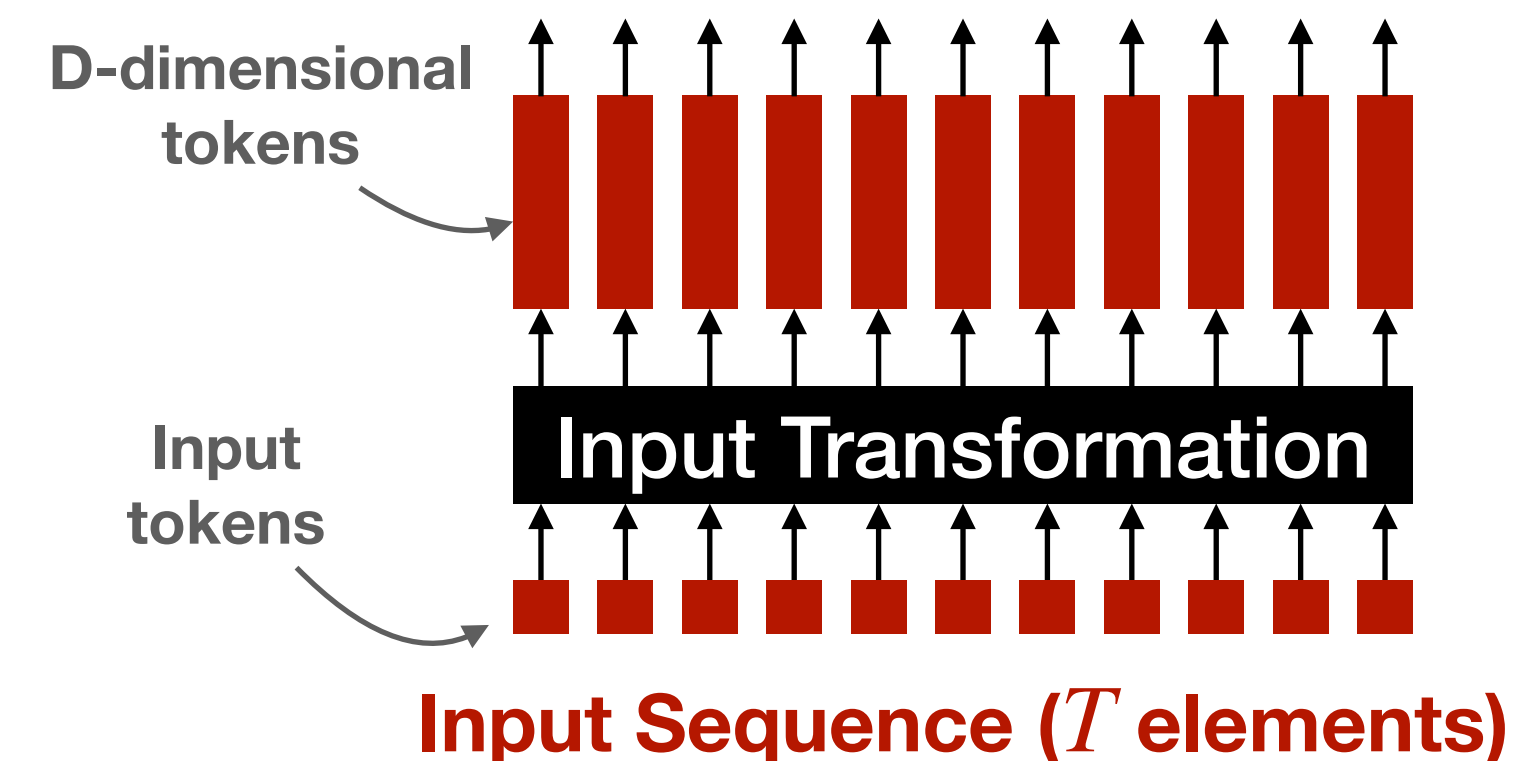
**Transformer** is a neural network  $f$  that iteratively transforms a sequence to another sequence and mixes the information between the sequence elements via **self-attention**



# Overview of Transformer Architecture

**Input transformation:** converts the input sequence elements into real-valued vector representations (aka **tokens**):

- maps a one-hot word vector to a real-valued vector
- extracts an image patch and flattens it into a vector




# Overview of Transformer Architecture

**Input transformation:** converts the input sequence elements into real-valued vector representations (aka **tokens**):

- maps a one-hot word vector to a real-valued vector
- extracts an image patch and flattens it into a vector

**Transformer block:** transforms a sequence of  $T$  vectors of dimension  $D$  into a new sequence of  $T$  vectors of dimension  $D$  using **self-attention** and **MLP sub-blocks**



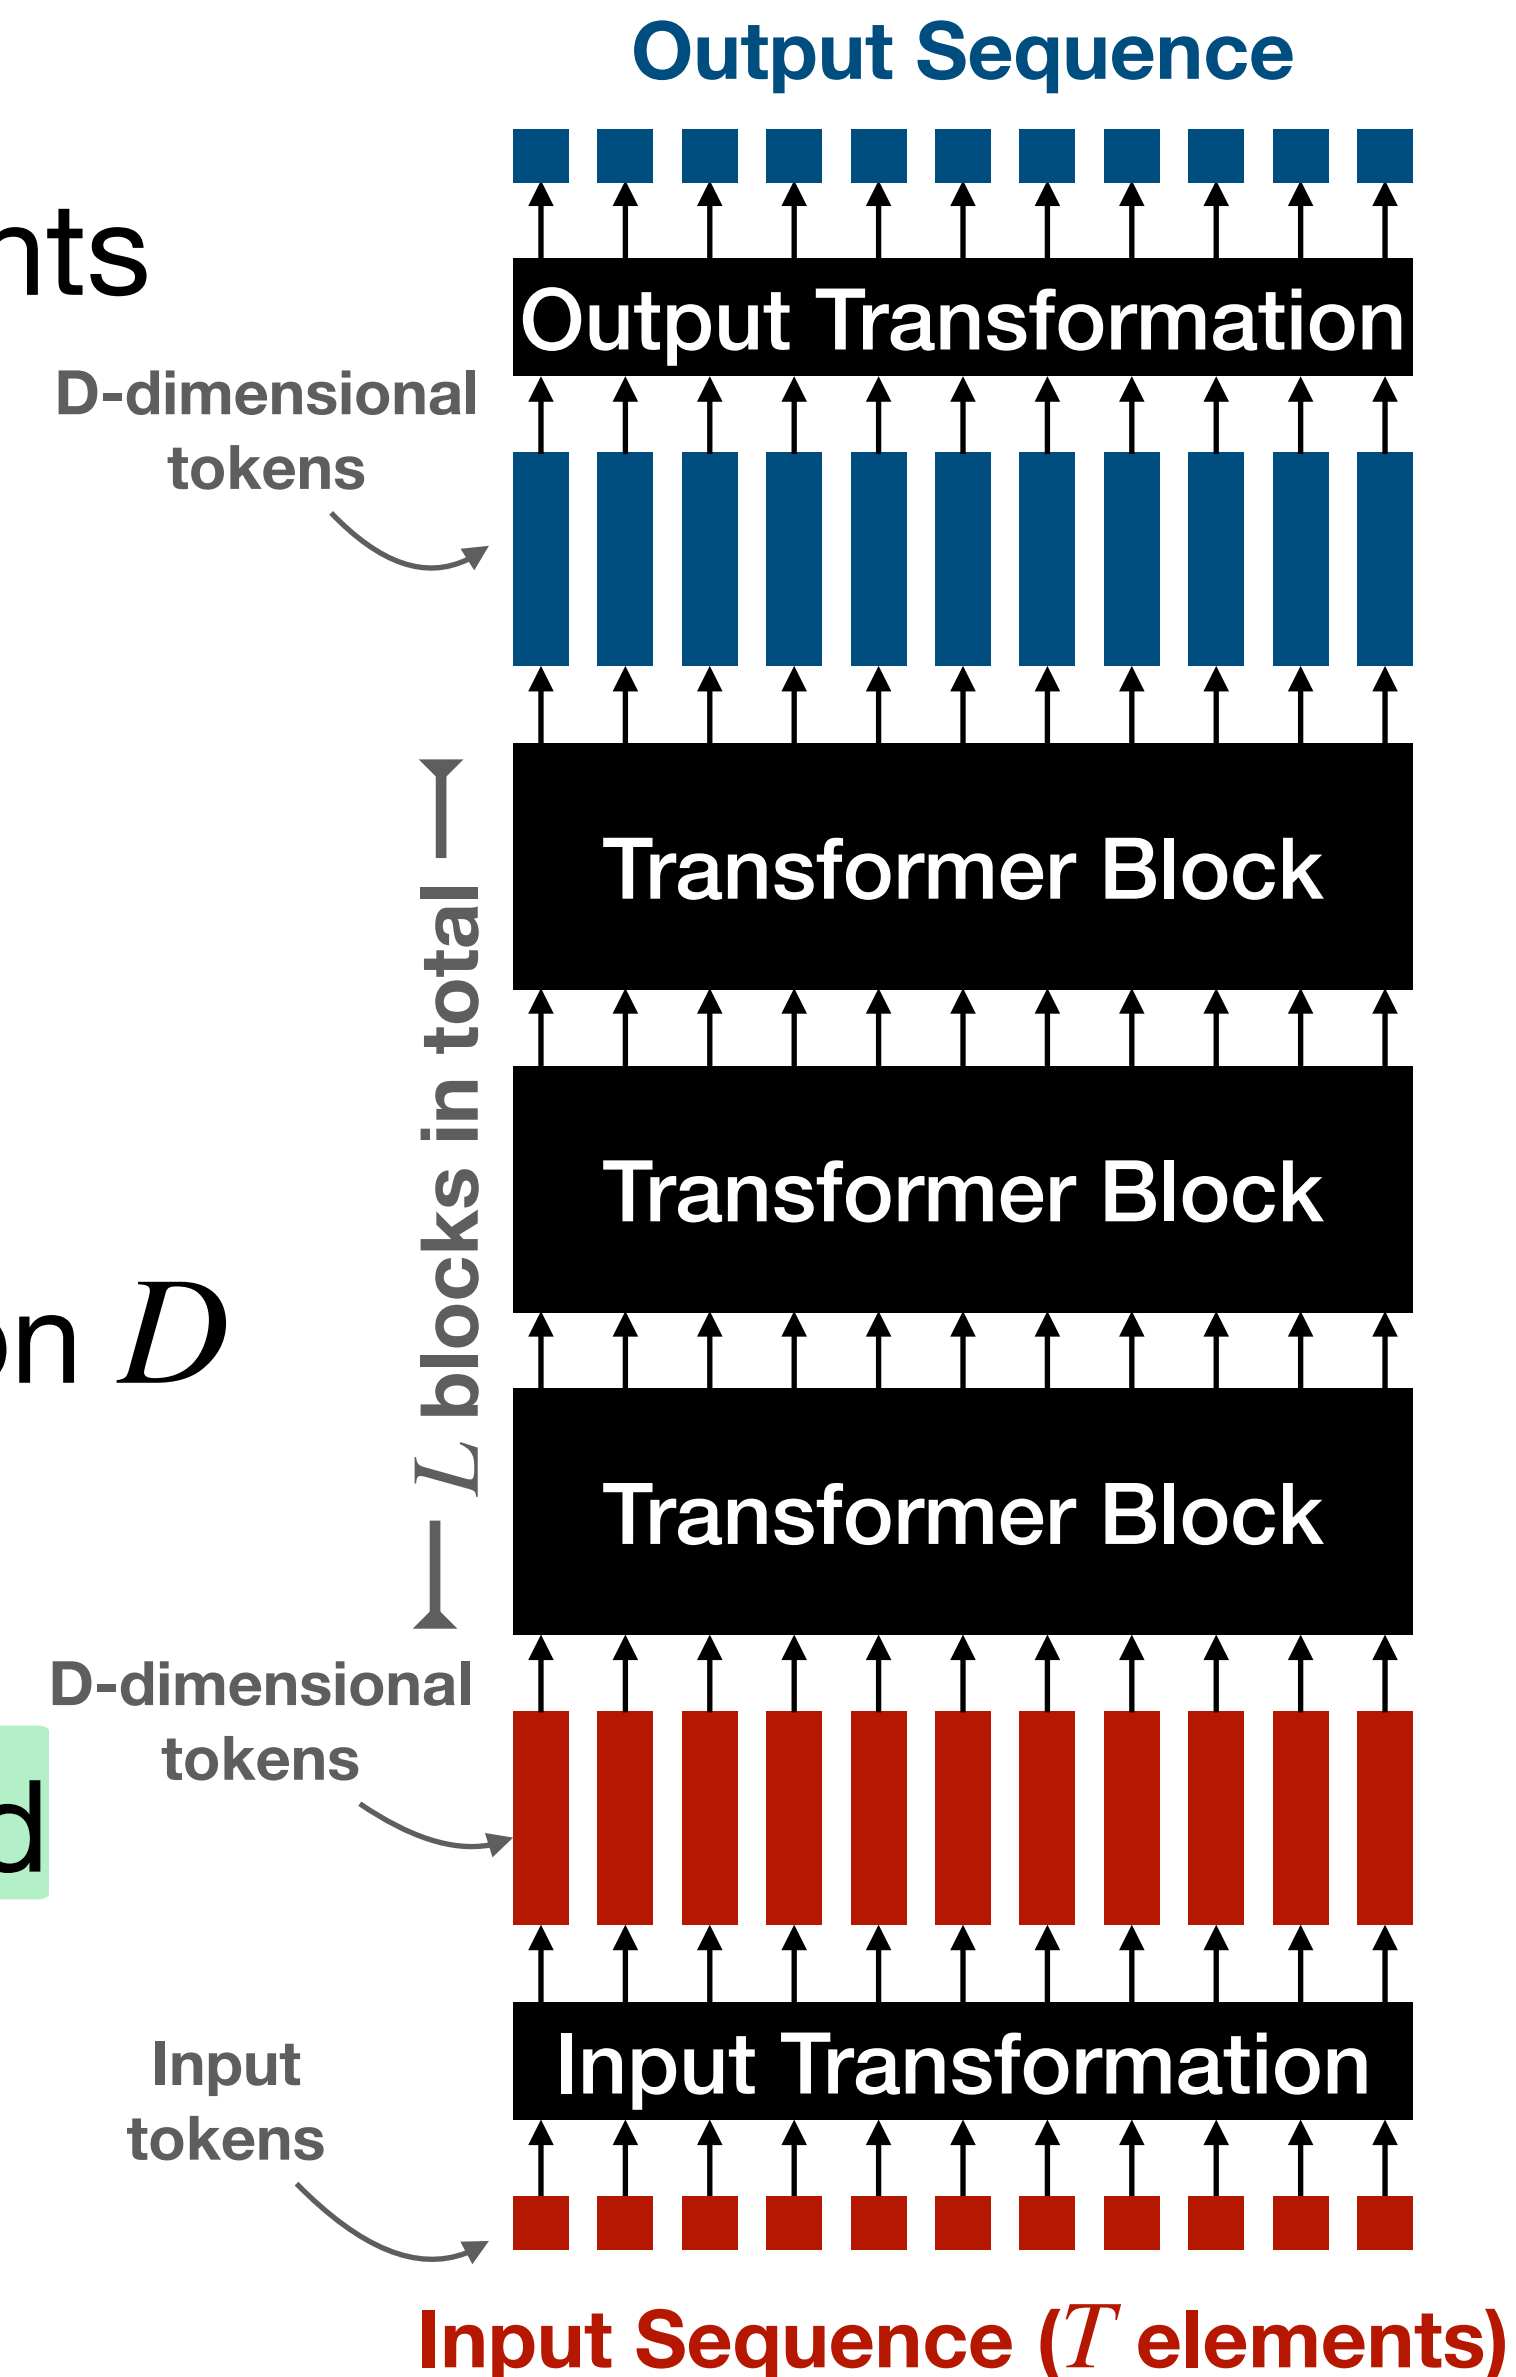
# Overview of Transformer Architecture

**Input transformation:** converts the input sequence elements into real-valued vector representations (aka **tokens**):

- maps a one-hot word vector to a real-valued vector
- extracts an image patch and flattens it into a vector

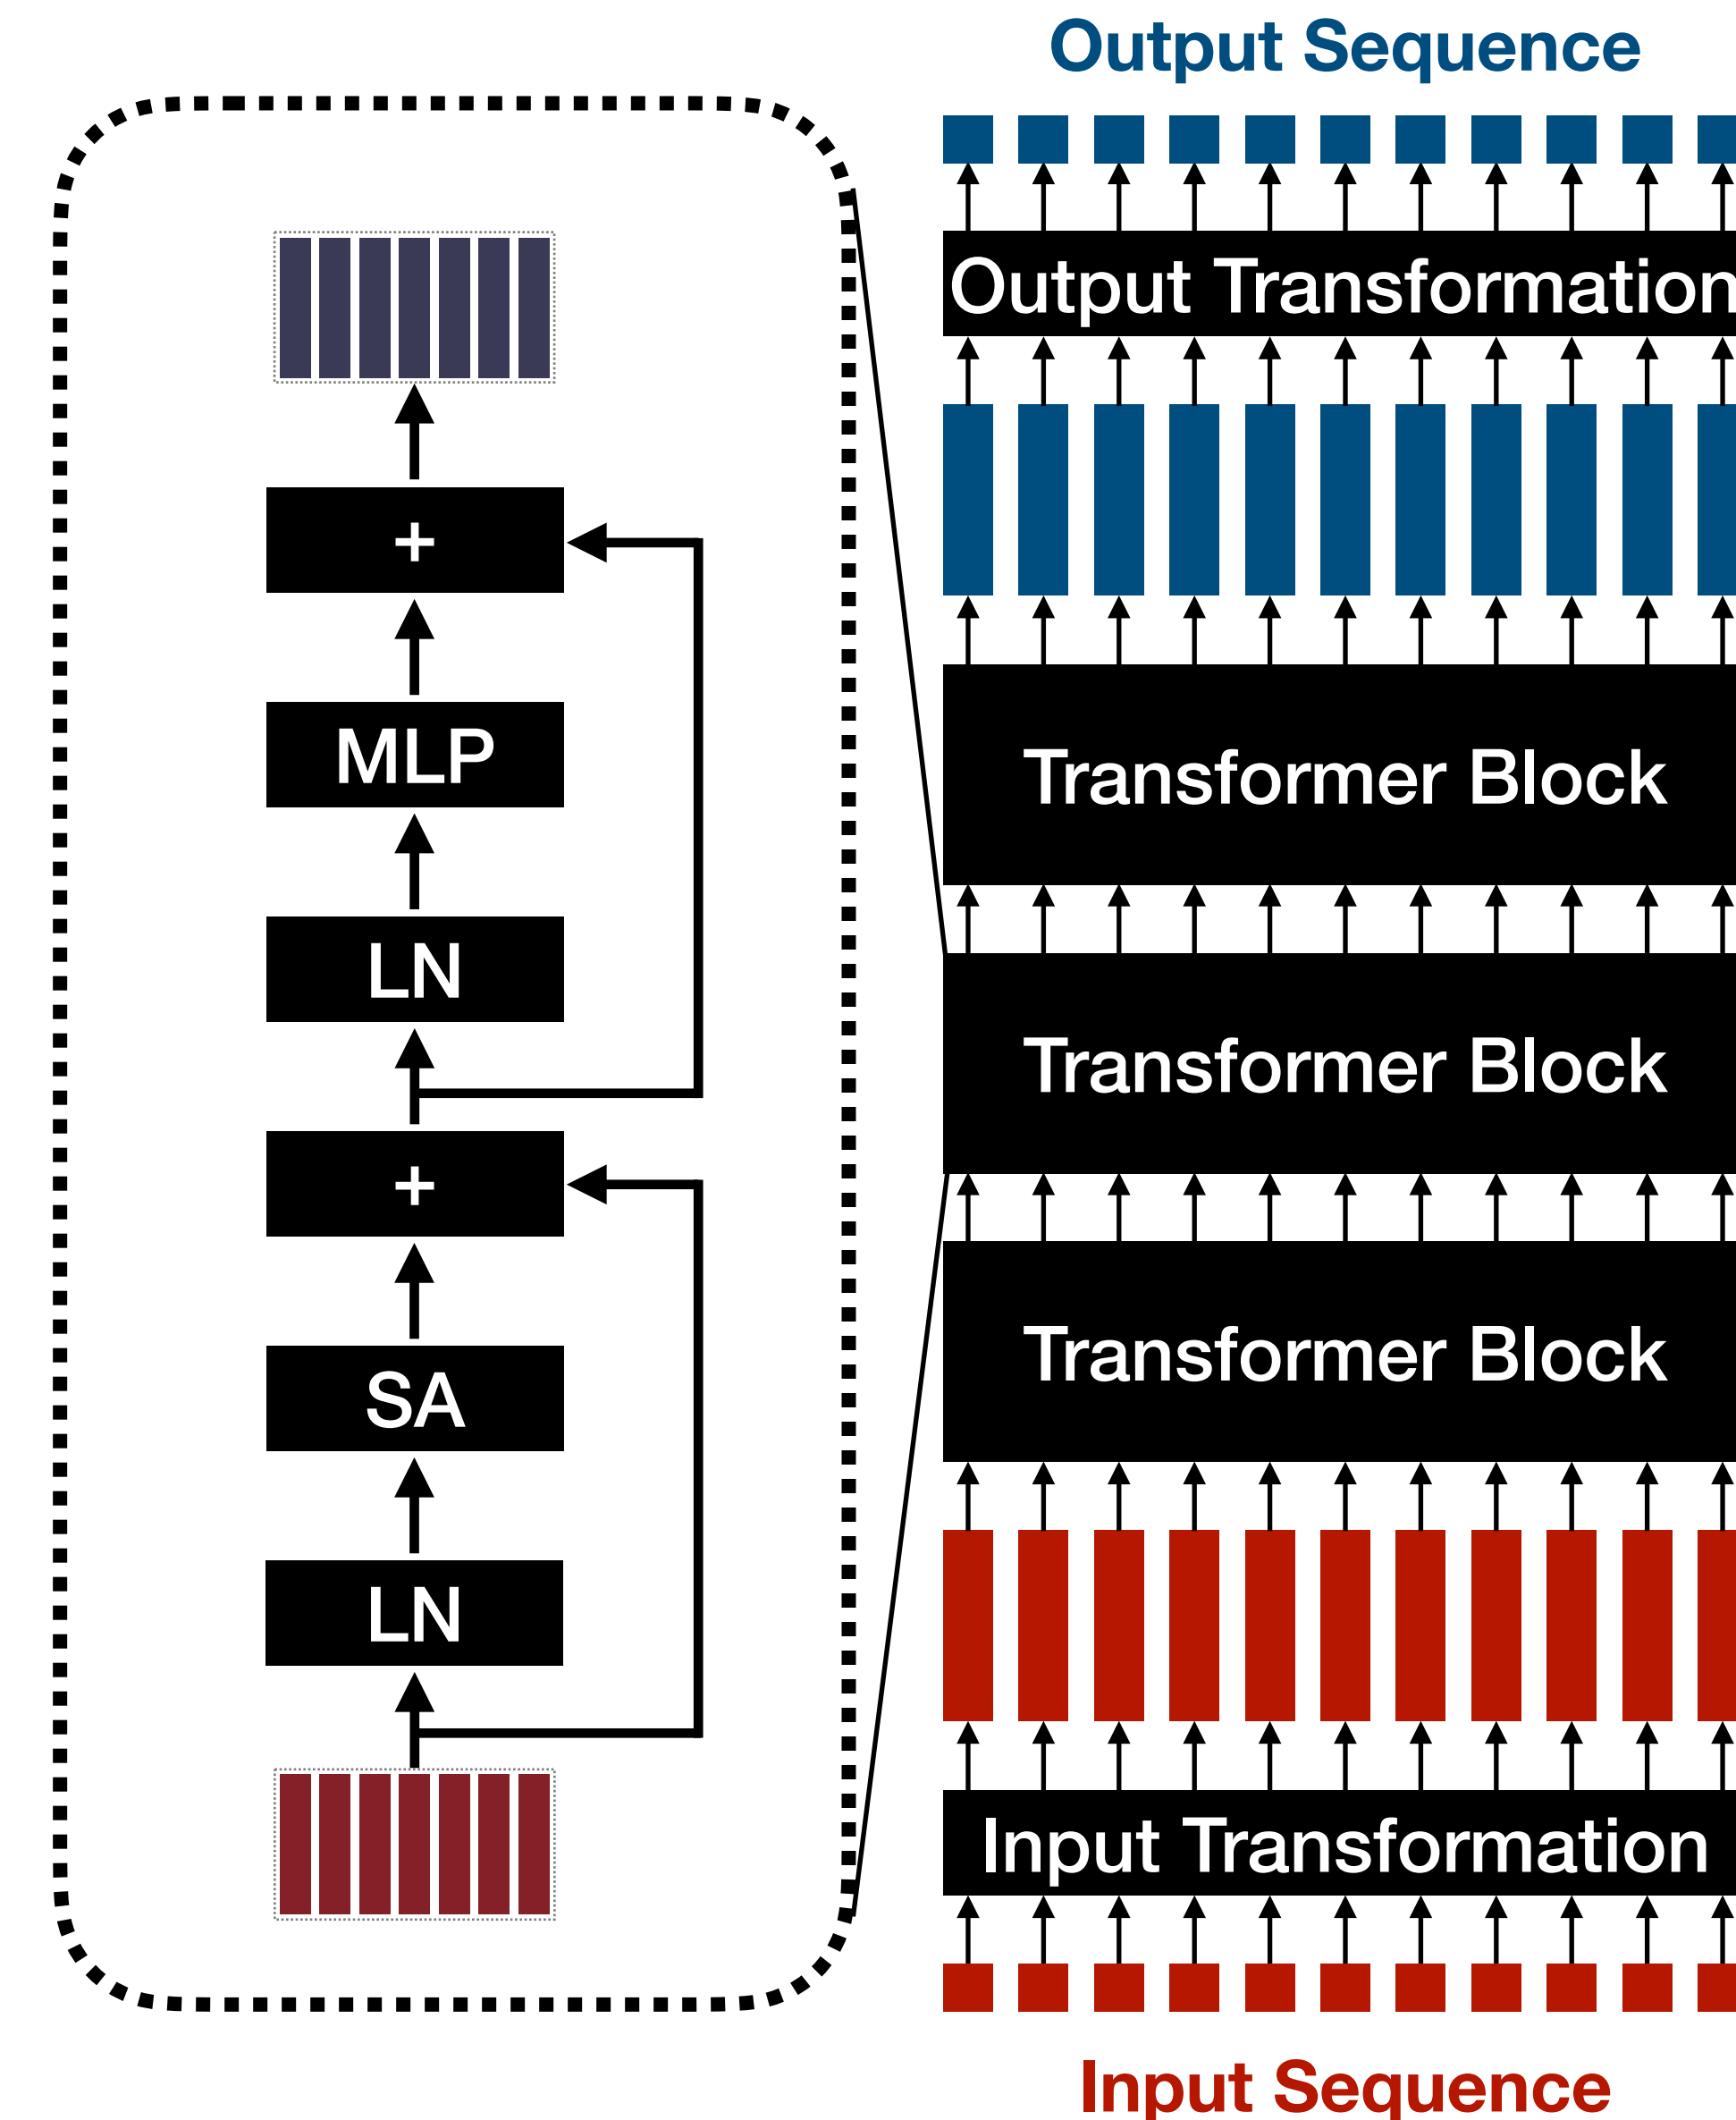
**Transformer block:** transforms a sequence of  $T$  vectors of dimension  $D$  into a new sequence of  $T$  vectors of dimension  $D$  using **self-attention** and **MLP sub-blocks**

**Output transformation:** converts the vectors to the desired output format (e.g., single-element sequence for classification, multiple-element sequence of words)



# Transformer Block

- **Self-Attention (SA):** mixes information **between** tokens
- **Multi-Layer Perceptron (MLP):** mixes information **within** each token
- Other standard components:
  - Skip connections are widely used
  - Layer normalization (LN) is usually placed at the start of a residual branch



# Input Transformations



# Text Token Embeddings

**Tokenization:** split the input text into a sequence of *input tokens* (typically word fragments + some special symbols) according to some predefined *tokenizer procedure*:

- Text: "<User:>Transformers are awesome!"
- Tokens: [<User token>, "Trans", "form", "ers\_", "are\_", "awe", "some", "!"]
- Token IDs: [0, 5124, 1029, 645, 3001, 6931, 7330, 10] (each token corresponds to some number  $i \in \{1, \dots, N_{vocab}\}$ )

**Token embedding:** maps each token ID  $i \in \{1, \dots, N_{vocab}\}$  into a real-valued vector  $\mathbf{w}_i \in \mathbb{R}^D$ :

- Token embeddings:  $[w_0, w_{5124}, w_{1029}, w_{645}, w_{3001}, w_{6931}, w_{7330}, w_{10}]$

➡ The whole input sequence of  $T$  tokens leads to an input matrix  $X = \begin{bmatrix} w_0^\top \\ w_{5124}^\top \\ \vdots \\ w_{10}^\top \end{bmatrix} \in \mathbb{R}^{T \times D}$

# Text Token Embeddings - Learning

- The matrix  $\mathbf{W}_{\text{emb}} = \begin{bmatrix} \mathbf{w}_1^\top \\ \vdots \\ \mathbf{w}_{N_{\text{vocab}}}^\top \end{bmatrix} \in \mathbb{R}^{N_{\text{vocab}} \times D}$  is learned via backpropagation, along with all other transformer parameters


- This can be seen as a matrix multiplication:

$$\mathbf{X} = \begin{bmatrix} \mathbf{e}_{i_1}^\top \\ \vdots \\ \mathbf{e}_{i_T}^\top \end{bmatrix} \mathbf{W}_{\text{emb}} \quad (\text{since } \mathbf{e}_i^\top \mathbf{W}_{\text{emb}} = (\mathbf{W}_{\text{emb}})_{i,:} = \mathbf{w}_i^\top)$$

- The tokenizer procedure is typically fixed in advance and not learned

# Image Patch Embeddings

- Divide image into patches of a given size (typical choice:  $16 \times 16$  pixels each)
- Flatten each patch into a vector of size  $16 \cdot 16 \cdot 3 = 768$  (height\*width\*color channels)
- Multiply each resulting vector by an embedding matrix  $\mathbf{W}_{\text{emb}} \in \mathbb{R}^{D \times 768}$  which is shared for all inputs
- Learn  $\mathbf{W}_{\text{emb}}$  through backpropagation, along with all other transformer parameters
- The whole input sequence of  $T$  embedded patches leads to an input matrix  $X \in \mathbb{R}^{T \times D}$



# Self-attention

# What is Self-attention?

$$A : \textit{tokens} \rightarrow \textit{tokens}$$

(using a weighted average)

Reminder: a token is simply a real-valued vector

**Self-attention** is a function that transforms a sequence of tokens to a new sequence of tokens using a **learned input-dependent weighted average**



# Self-Attention

Define  $K, Q, V$  from the **same** input sequence  $X \in \mathbb{R}^{T \times D}$


- **Keys:**  $K = XW_K \in \mathbb{R}^{T \times D_K}$
  - **Queries:**  $Q = XW_Q \in \mathbb{R}^{T \times D_K}$
  - **Values:**  $V = XW_V \in \mathbb{R}^{T \times D_V}$
- $W_K, W_Q \in \mathbb{R}^{T \times D_K}, W_V \in \mathbb{R}^{T \times D_V}$  are parameters

The output of self-attention is then given by:

$$Z = \text{softmax} \left( \frac{QK^\top}{\sqrt{D_K}} \right) V$$

→  $\text{softmax}(\cdot)$  is applied row-wise

→ Quadratic computational complexity  $O(T^2)$



# Attention as a Weighted Average

- $T$  input and output tokens:  $V \in \mathbb{R}^{T \times D_V}, Z \in \mathbb{R}^{T \times D}$

**Notation:** throughout this lecture, the  $j$ -th rows of  $V$  and  $Z$  are denoted by  $\mathbf{v}_j$  and  $\mathbf{z}_j$


- Outputs are a **weighted average** of the inputs:

$$\mathbf{z}_i = \sum_{j=1}^T p_{i,j} \mathbf{v}_j \quad \text{or in matrix form } Z = PV$$



- Weighting coefficients  $P \in [0,1]^{T \times T}$  form valid probability distributions over the input tokens:

$$\Rightarrow \sum_{j=1}^T p_{i,j} = 1 \quad (\text{i.e., each row sums to one})$$

Input tokens  $V$



Output tokens  $Z$




# The Weighting Coefficients $P$

- **Query tokens**  $Q \in \mathbb{R}^{T \times D_K}$  (one query per output token)
- **Key tokens**  $K \in \mathbb{R}^{T \times D_K}$  (one key per input token)
- Determine weight  $p_{i,j}$  based on **how similar  $q_i$  and  $k_j$  are**
  - Use inner product to obtain raw similarity scores
  - Normalize with softmax (scaled the temperature by  $\sqrt{D_K}$ ) to obtain a probability distribution
- This can be expressed as:

**Element-wise:** 
$$p_{i,j} = \frac{\exp\left(\mathbf{q}_i \mathbf{k}_j^\top / \sqrt{D_K}\right)}{\sum_{t=1}^T \exp\left(\mathbf{q}_i \mathbf{k}_t^\top / \sqrt{D_K}\right)}$$

**Matrix form:** 
$$P = \text{softmax}\left(\frac{QK^\top}{\sqrt{D_K}}\right)$$

← The softmax is applied on each row *independently*



**Computation complexity:**  
 $O(T \times T)$

# The Weighting Coefficients $P$

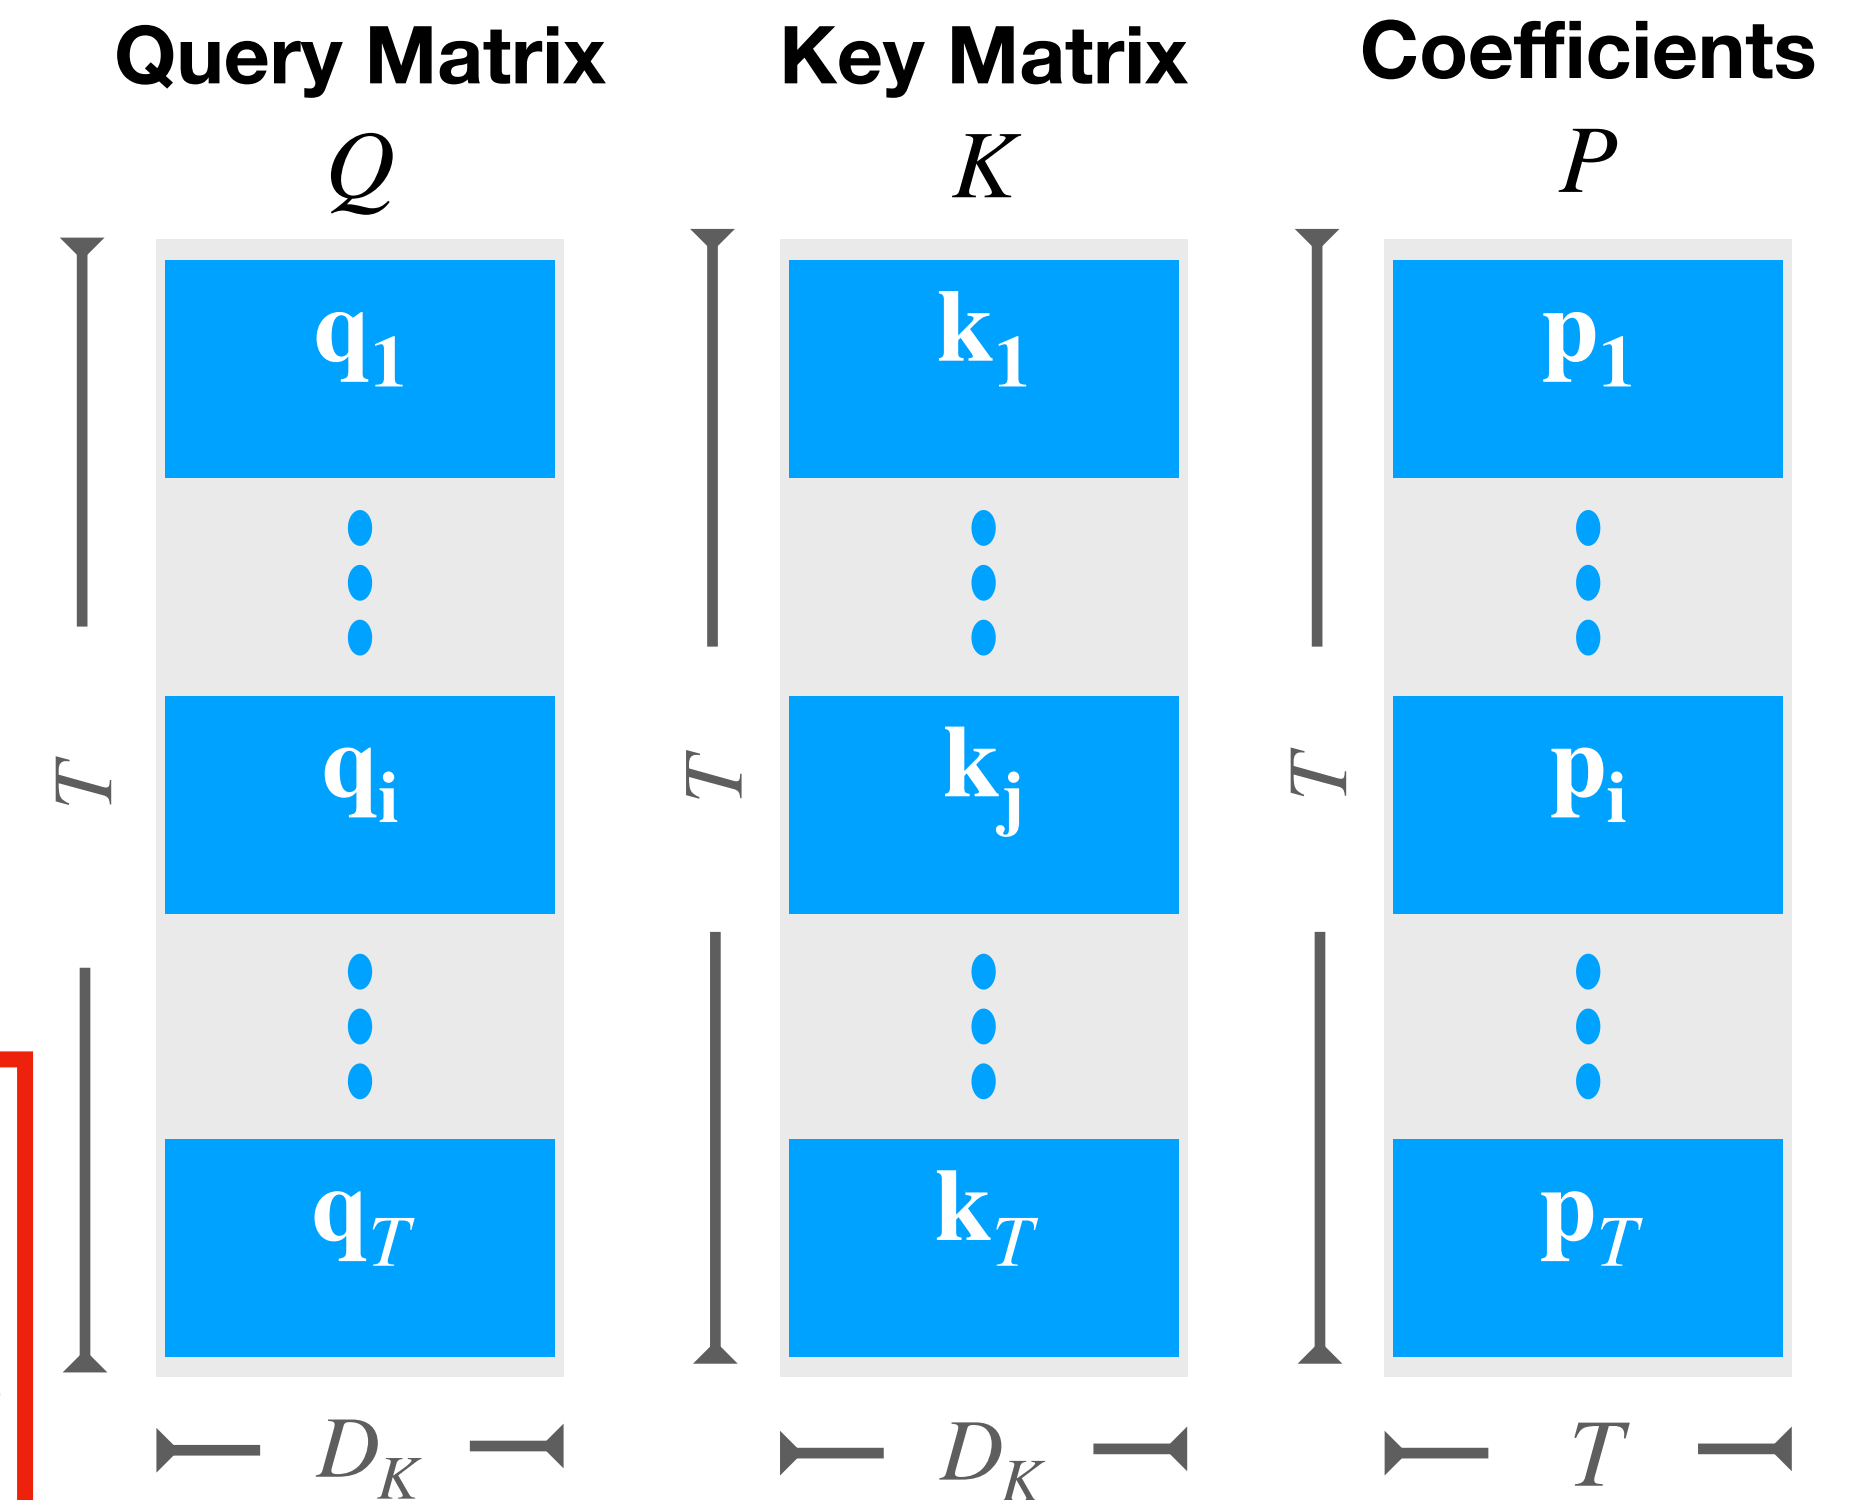
- **Query tokens**  $Q \in \mathbb{R}^{T \times D_K}$  (one query per output token)
- **Key tokens**  $K \in \mathbb{R}^{T \times D_K}$  (one key per input token)
- Determine weight  $p_{i,j}$  based on **how similar  $q_i$  and  $k_j$**  are
  - Use inner product to obtain raw similarity scores
  - Normalize with softmax (scaled the temperature by  $\sqrt{D_K}$ ) to obtain a probability distribution

In some applications, **causal masking** is used:

**Sum until position  $i$ :**  $p_{i,j} = \frac{\exp(q_i k_j^\top / \sqrt{D_K})}{\sum_{t=1}^i \exp(q_i k_t^\top / \sqrt{D_K})}$  for  $j \leq i$  and  $p_{i,j} = 0$  otherwise

**Mask before softmax:**  $P = \text{softmax} \left( \textcolor{red}{M} + \frac{QK^\top}{\sqrt{D_K}} \right)$

where  $M \in \mathbb{R}^{T \times T}$  is the matrix  $M_{ij} = -\infty$  for  $j > i$  and  $M_{i,j} = 0$  otherwise



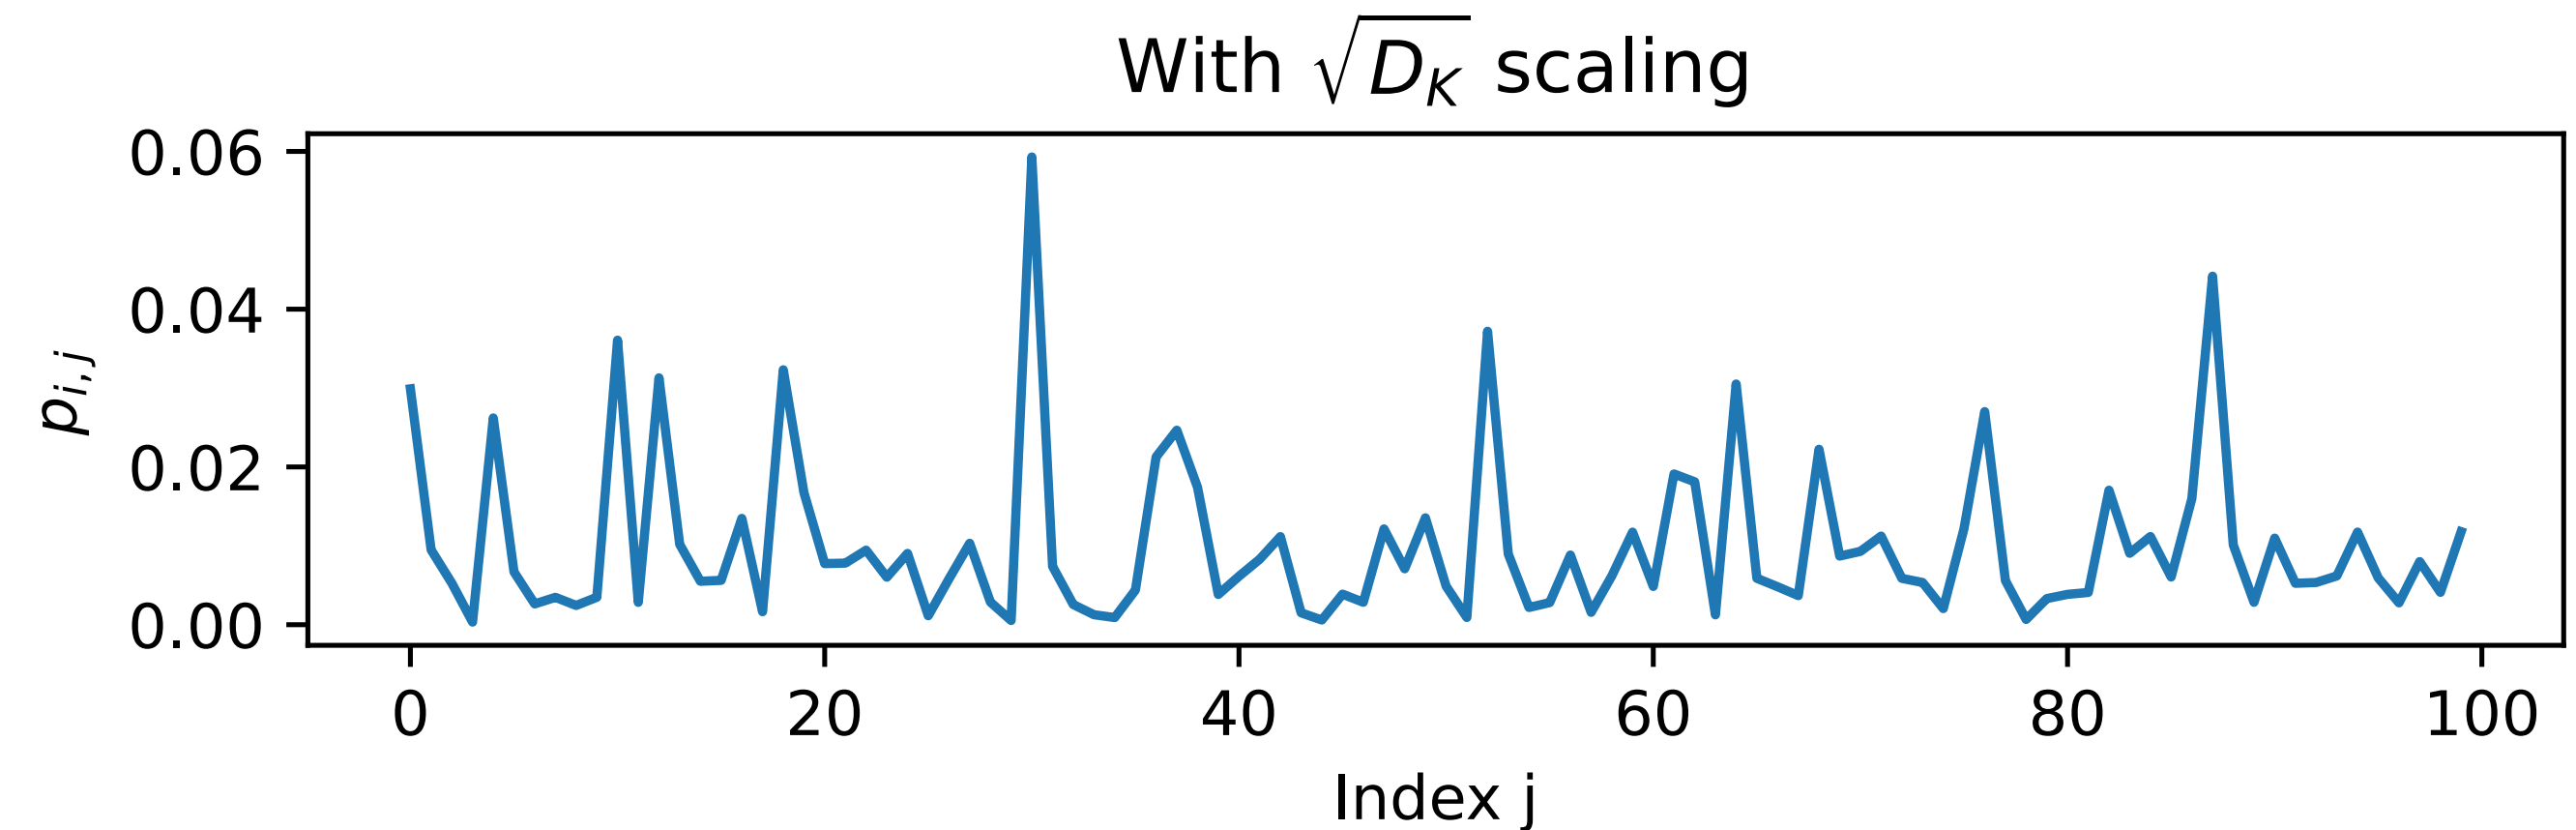
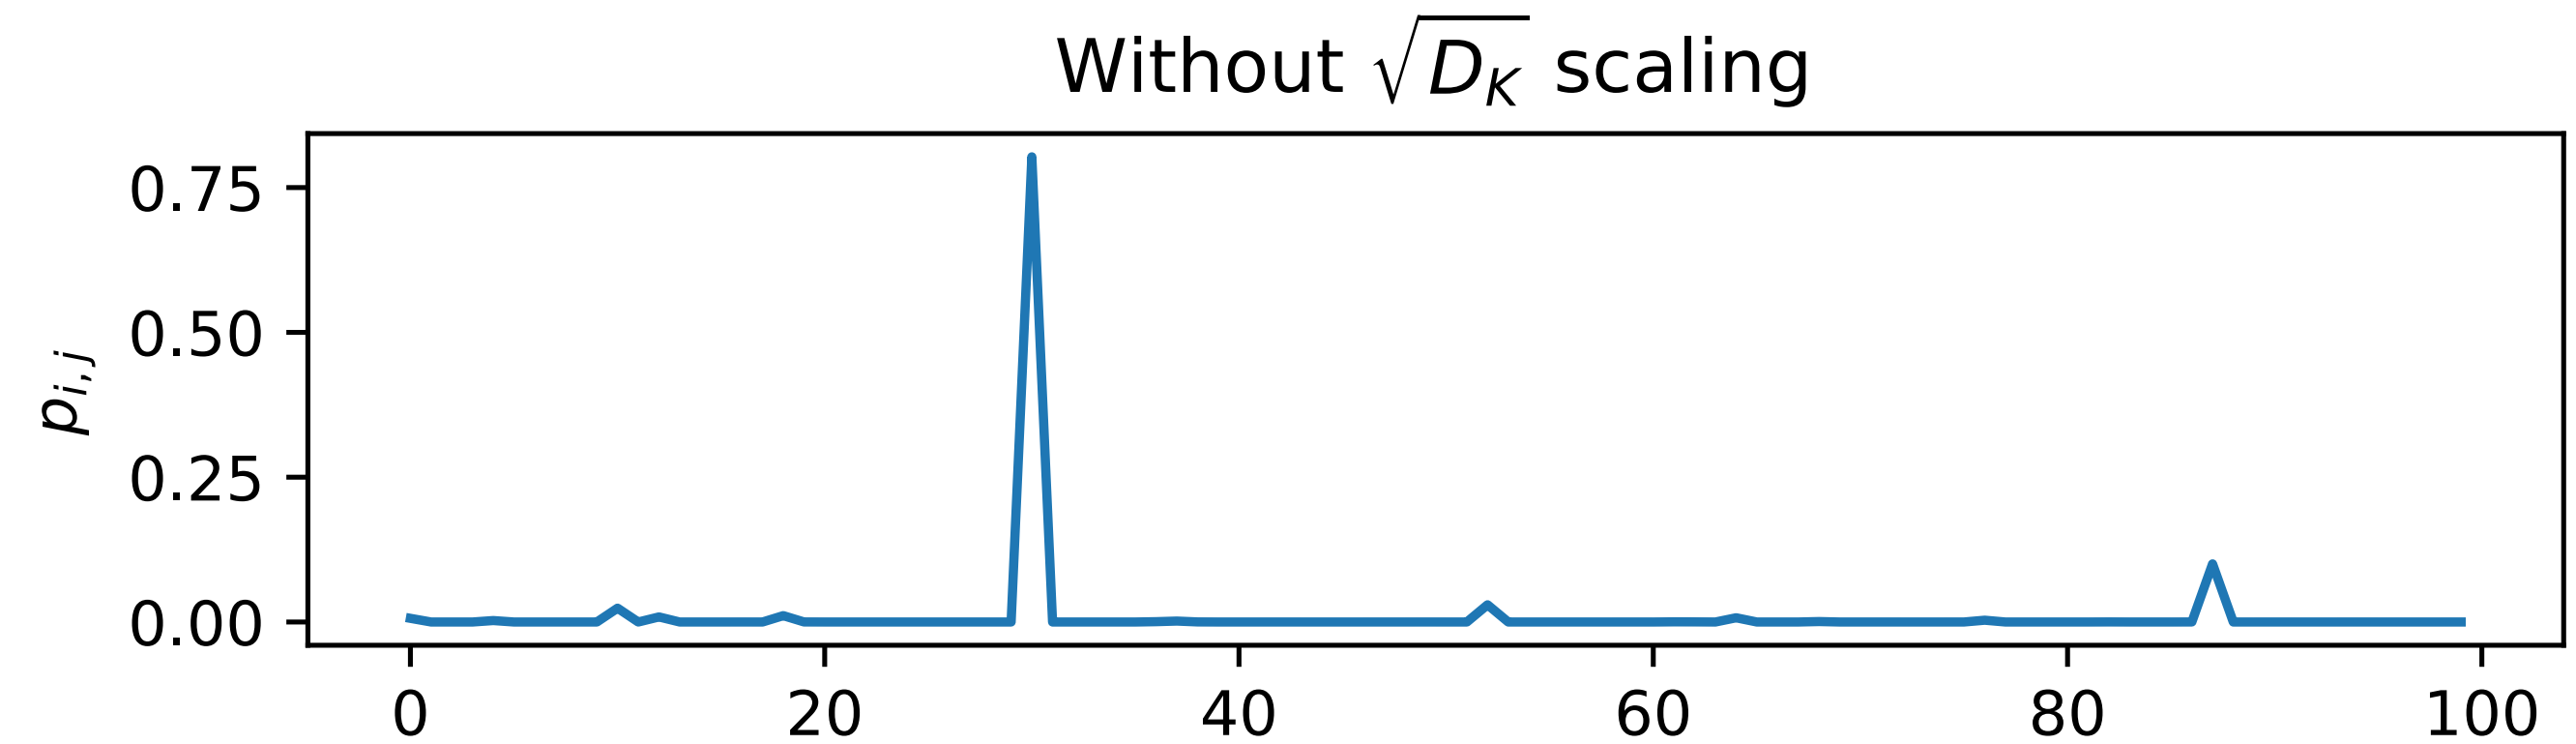
**Computation complexity:**

$$O(T \times T)$$

# Why Use the $1/\sqrt{D_K}$ Scaling?

$$P = \text{softmax} \left( \frac{QK^\top}{\sqrt{D_K}} \right)$$

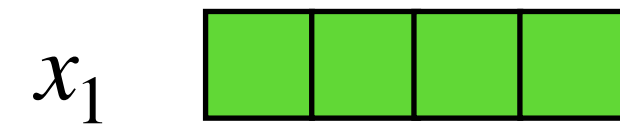
- **Without scaling:** sharp distribution of the attention weights  $p_{i,j}$  at random initialization
- The model takes much more time to adjust from the initial peak due to vanishing gradients
- The  $1/\sqrt{D_K}$  scaling ensures uniformity at initialization and faster convergence




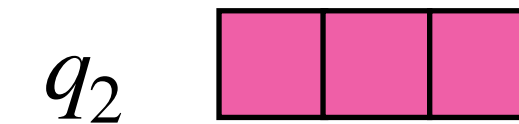
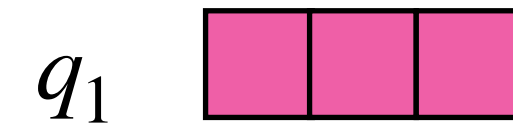


# Self-Attention: Step-by-Step


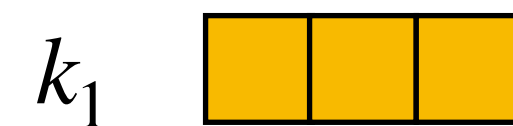
Input




Queries



Keys



Values

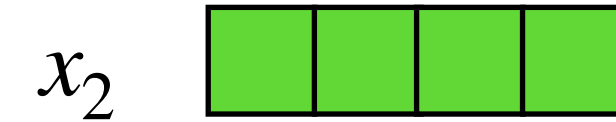
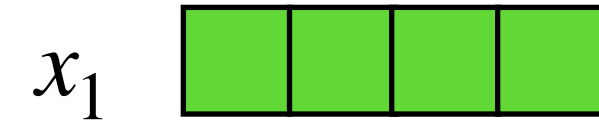


$$\begin{aligned} Q &= XW_Q \\ K &= XW_K \\ V &= XW_V \end{aligned}$$

Multiplying the input by the Q/K/V weight matrices, we create a query, a key and a value projection of each input of the input sequence

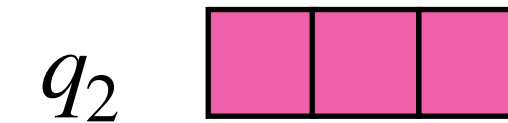
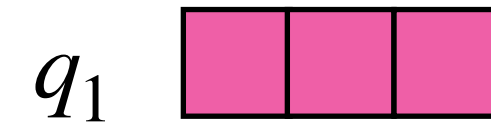
# Self-Attention: Step-by-Step

Input

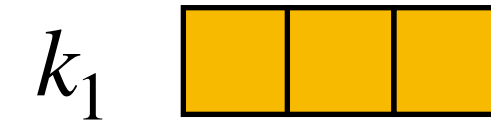


Step 1: create query, key and value vectors  
for each input token

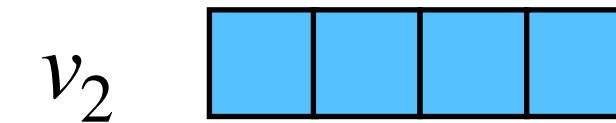
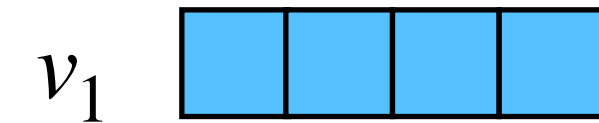
Queries



Keys



Values



$$\begin{aligned} Q &= XW_Q \\ K &= XW_K \\ V &= XW_V \end{aligned}$$

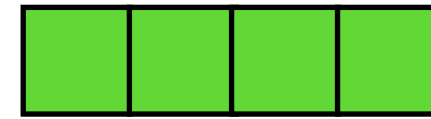
# Self-Attention: Step-by-Step

Input

$x_1$

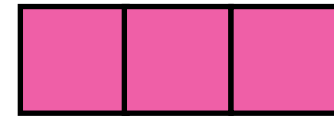


$x_2$

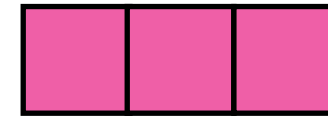


Queries

$q_1$



$q_2$



Keys

$k_1$

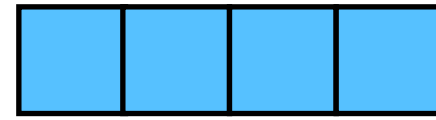


$k_2$

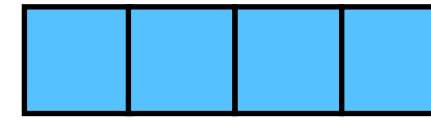


Values

$v_1$



$v_2$



Score

$$q_1 k_1^T = 102$$

$$q_1 k_2^T = 99$$

Step 2: calculate the scores by taking scalar product of the query and key vectors

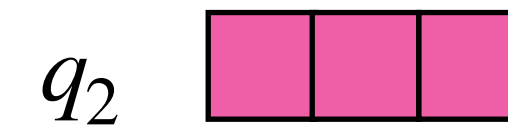
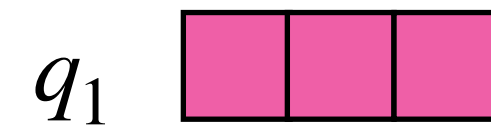
$$QK^T = XW_QW_K^TX^T$$

# Self-Attention: Step-by-Step

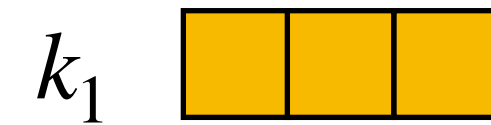
Input



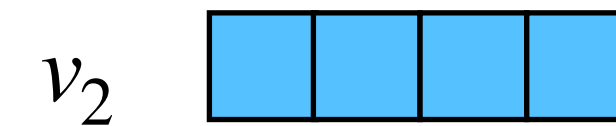
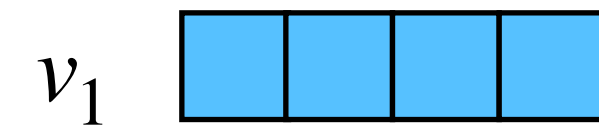
Queries



Keys



Values



Score

$$q_1 k_1^\top = 102$$

$$q_1 k_2^\top = 99$$

Divide by  $\sqrt{D_K}$

$$\frac{q_1 k_1^\top}{\sqrt{D_K}} = 58.9$$

$$\frac{q_1 k_2^\top}{\sqrt{D_K}} = 57.2$$

Softmax

$$p_{1,1} = 0.85$$

$$p_{1,2} = 0.15$$

Step 3: divide the scores by  $\sqrt{D_K}$

Step 4: Compute the softmax of these values

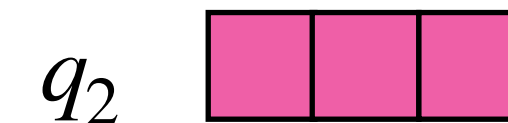
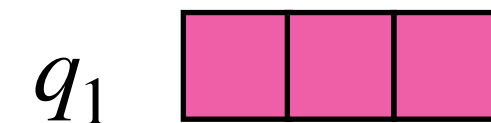
$$P = \text{softmax} \left( \frac{QK^\top}{\sqrt{D_K}} \right)$$

# Self-Attention: Step-by-Step

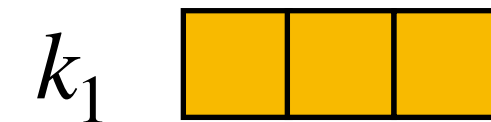
Input



Queries



Keys



Values



Score

$$q_1 k_1^T = 102$$

$$q_1 k_2^T = 99$$

Divide by  $\sqrt{D_K}$

$$\frac{q_1 k_1^T}{\sqrt{D_K}} = 58.9$$

$$\frac{q_1 k_2^T}{\sqrt{D_K}} = 57.2$$

Softmax

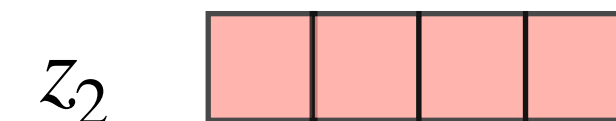
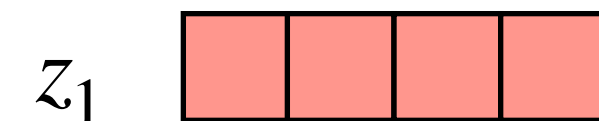
$$p_{1,1} = 0.85$$

$$p_{1,2} = 0.15$$

Softmax\*Value



Sum



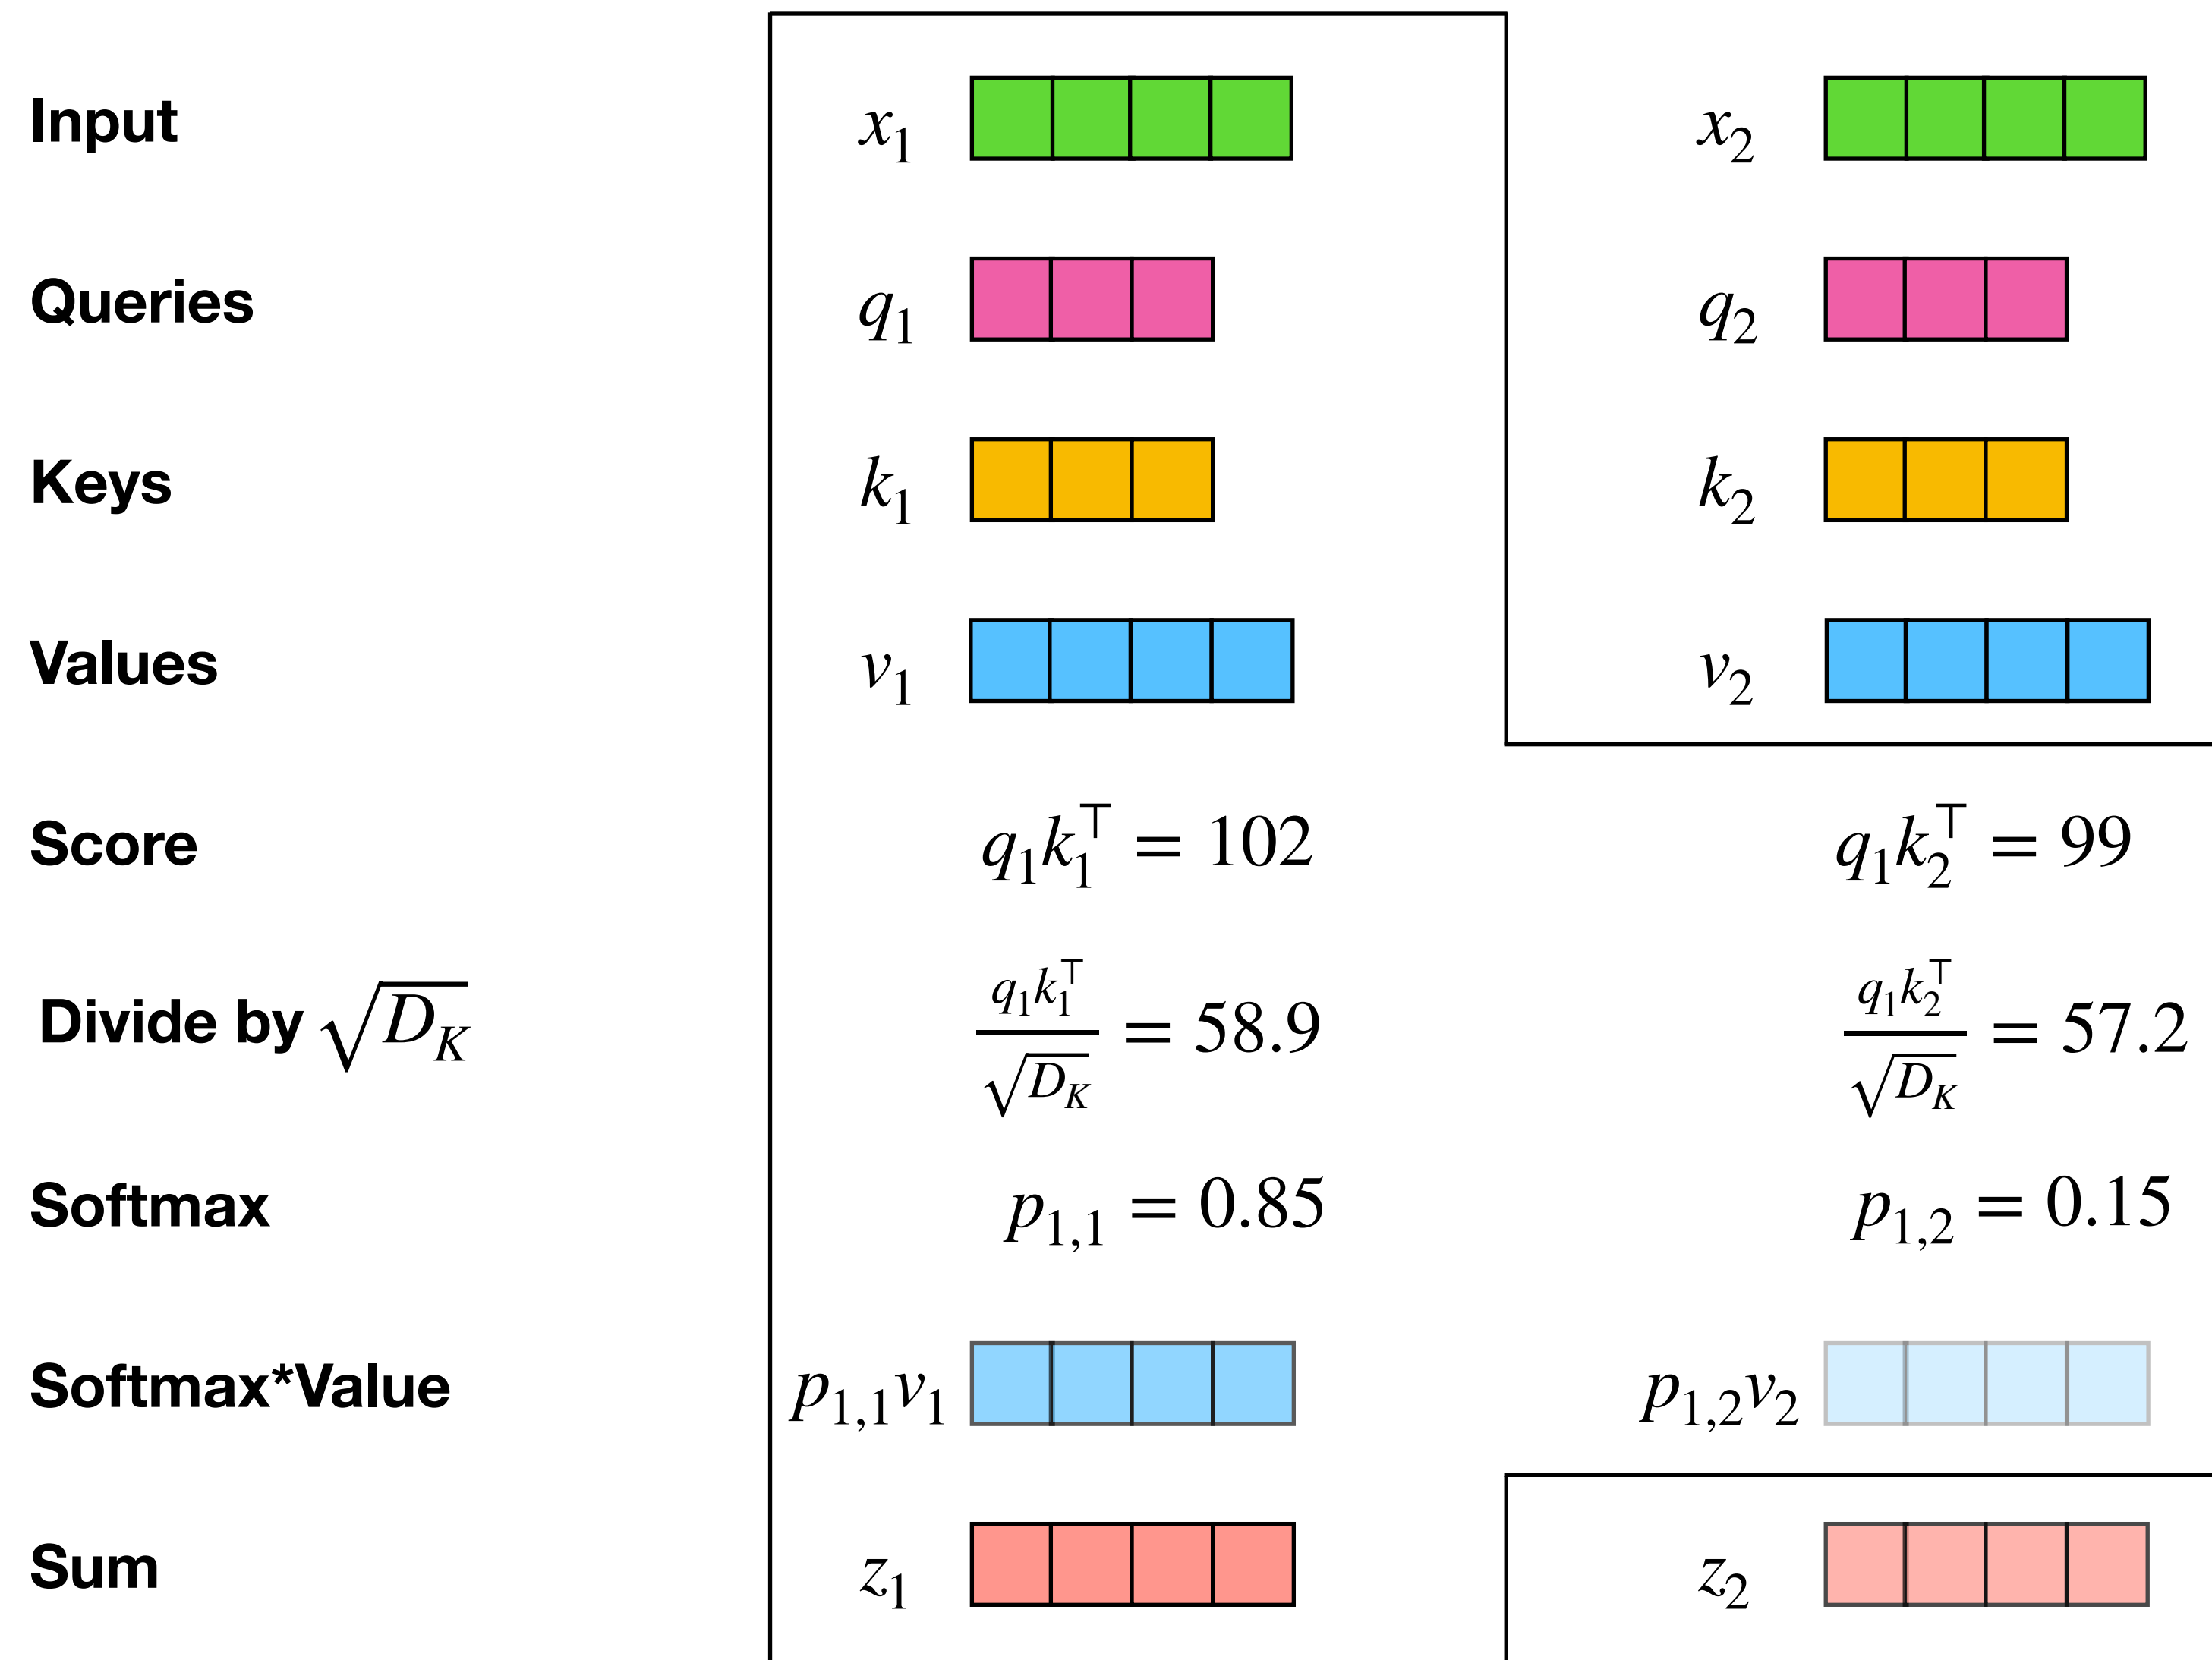
Step 5: Multiply each value vector  
by the softmax score

Step 6: Sum up the weighted value vectors

$$Z = PV$$



# Self-Attention: Step-by-Step



$$Z = \text{softmax} \left( \frac{XW_Q W_K^\top X^\top}{\sqrt{D_K}} \right) XW_V$$

# Multi-Head Self-Attention

- It is desirable to have multiple attention patterns per layer, similar to having multiple convolutions in a convolutional layer

➡ Run  $H$  Self-Attention “heads” in parallel

- The output of head  $h$  is given by:

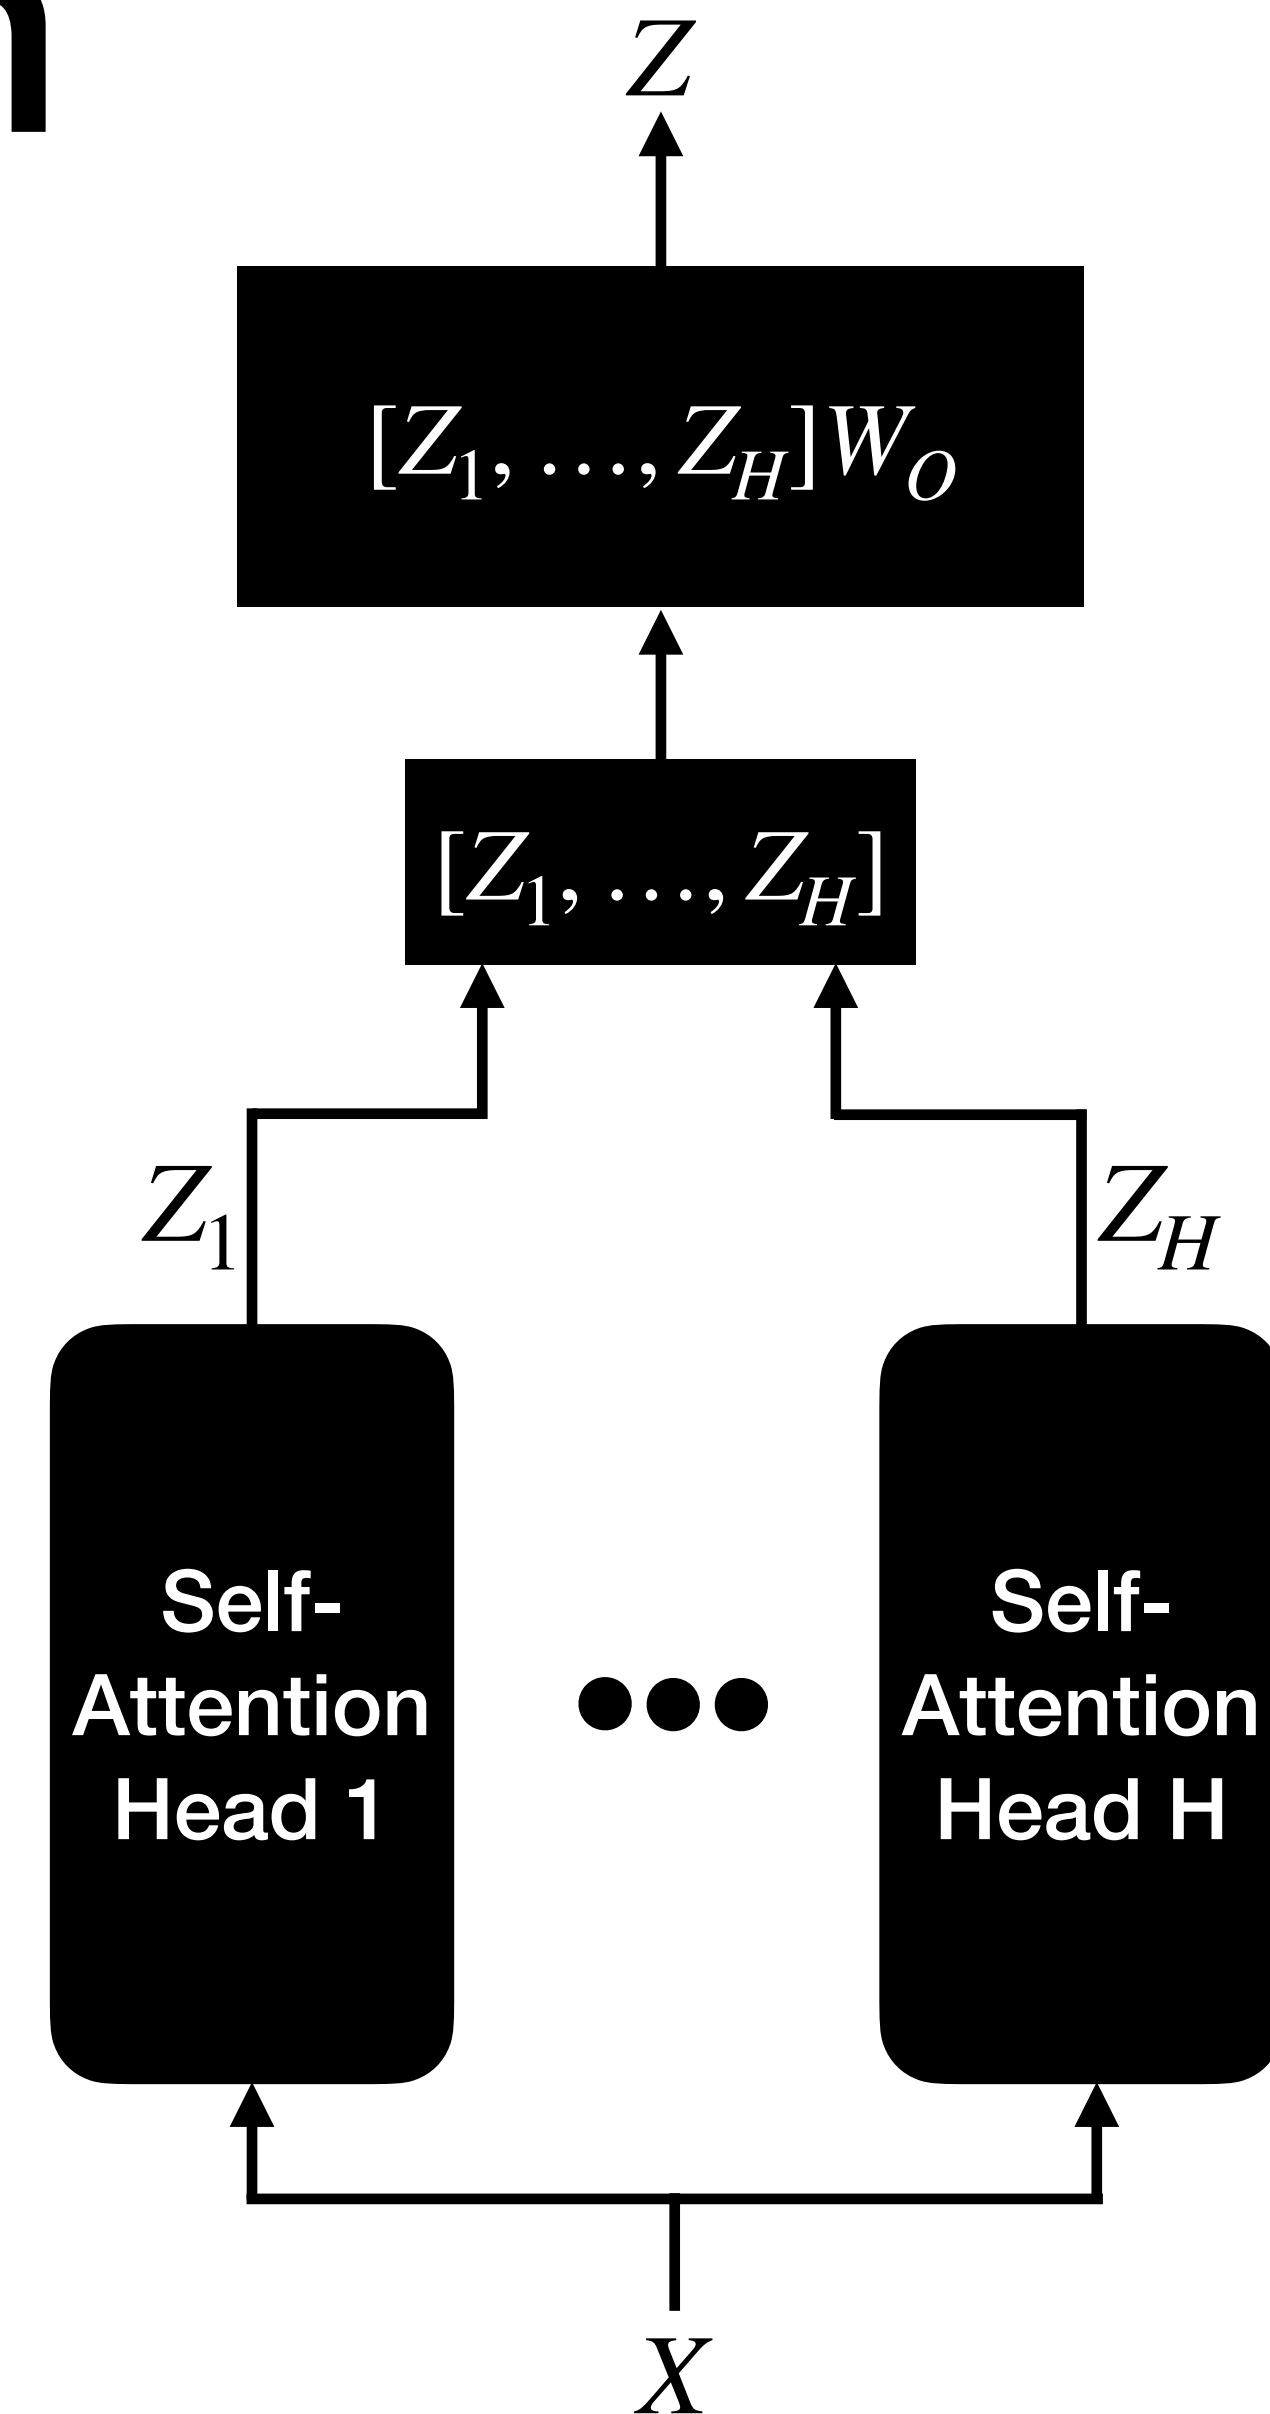
$$Z_h = \text{softmax} \left( \frac{XW_{Q,h}W_{K,h}^\top X^\top}{\sqrt{D_K}} \right) XW_{V,h}$$

$$W_{V,h} \in \mathbb{R}^{D \times D_V}, W_{K,h} \in \mathbb{R}^{D \times D_K}, W_{Q,h} \in \mathbb{R}^{D \times D_K}$$

- The final output is obtained by concatenating head-outputs and applying a linear transformation

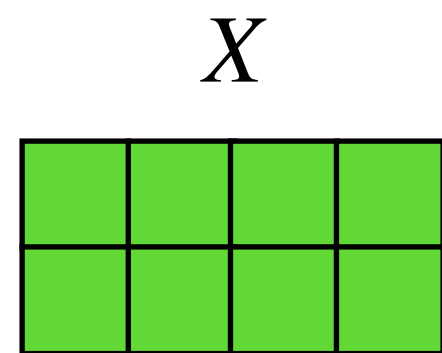
$$Z = [Z_1, \dots, Z_H]W_O$$

where  $W_O \in \mathbb{R}^{HD_V \times D}$  is learned via backpropagation



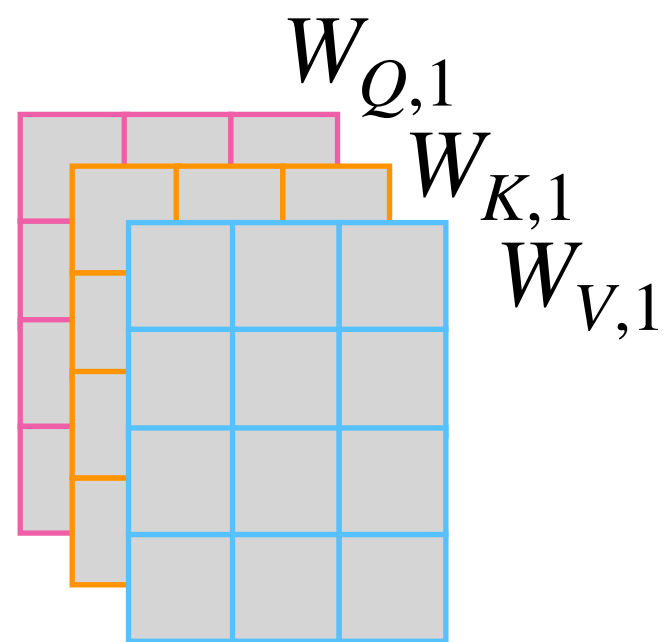
# Multi-Head Self-Attention: recap

1) Input

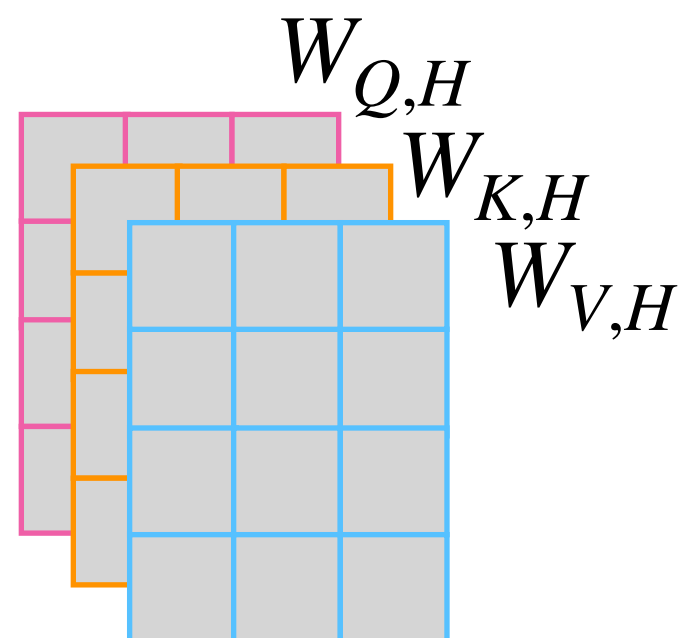


2) Split into  $H$  heads


We multiply  $X$  by weight matrices




...




3) Calculate attention using the resulting  $Q_h, K_h, V_h$  matrices



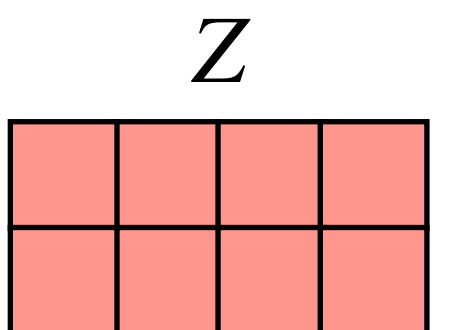
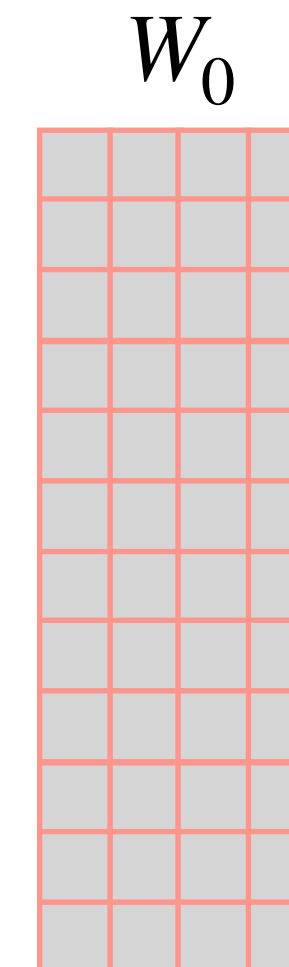

...



4) Concatenate the resulting matrices  $Z_h$  and multiply by  $W_0$  to obtain the final output  $Z$  of the self-attention layer



...



# Positional information

# Attention does not account for the order of input

For a permutation matrix  $R \in \{0,1\}^{T \times T}$  we have:

$$Z_R = \text{softmax} \left( \frac{RXW_Q W_K^T X^T R^T}{\sqrt{D_K}} \right) RXW_V$$

**Permute every X in original formula**

$$= R \text{softmax} \left( \frac{XW_Q W_K^T X^T R^T}{\sqrt{D_K}} \right) RXW_V$$

**Since softmax is computed row-wise**

$$= R \text{softmax} \left( \frac{XW_Q W_K^T X^T}{\sqrt{D_K}} \right) R^T RXW_V$$

**Reordering the terms in the softmax sum does not affect the output**

$$= RPR^{-1}RXW_V$$

**For a permutation matrix: transpose=inverse**

$$= RPXW_V$$

Which is equivalent to a permutation of the original output  $Z = PV$

# Positional Information in Transformers

- In practice, the input order matters:

*"She prefers cats to dogs"  $\neq$  "She prefers dogs to cats"*

- **Solution:** incorporate a positional encoding in the network which is a function from the position to a feature vector  $\text{pos} : \{1, \dots, T\} \rightarrow \mathbb{R}^D$
- **The most basic choice** is to add a positional embedding  $W_{\text{pos}}$  corresponding to each token's position  $t$  to the input embedding.  $W_{\text{pos}} \in \mathbb{R}^{T \times D}$  is learned via backpropagation along with the other parameters:

$$\mathbf{X} = \begin{bmatrix} \mathbf{e}_{i_1}^\top \\ \vdots \\ \mathbf{e}_{i_T}^\top \end{bmatrix} \mathbf{W}_{\text{emb}} + \begin{bmatrix} \mathbf{e}_1^\top \\ \vdots \\ \mathbf{e}_T^\top \end{bmatrix} \mathbf{W}_{\text{pos}}$$


- Numerous hand-crafted positional encodings exist (active area of research!)



**MLP**

# Mixing Information within Tokens

- **MLP** mixes information within each token
- Apply the same transformation to each token independently:
$$MLP(X) = \varphi(XW_1)W_2$$
- Matrices  $W_1, W_2 \in \mathbb{R}^{D \times D}$  learned via backprop
- Non-linearity  $\varphi$  in between (e.g., ReLU or GeLU)
- The model may also include learned bias terms

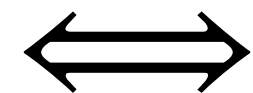


# Mixing Information within Tokens

- **MLP** mixes information within each token

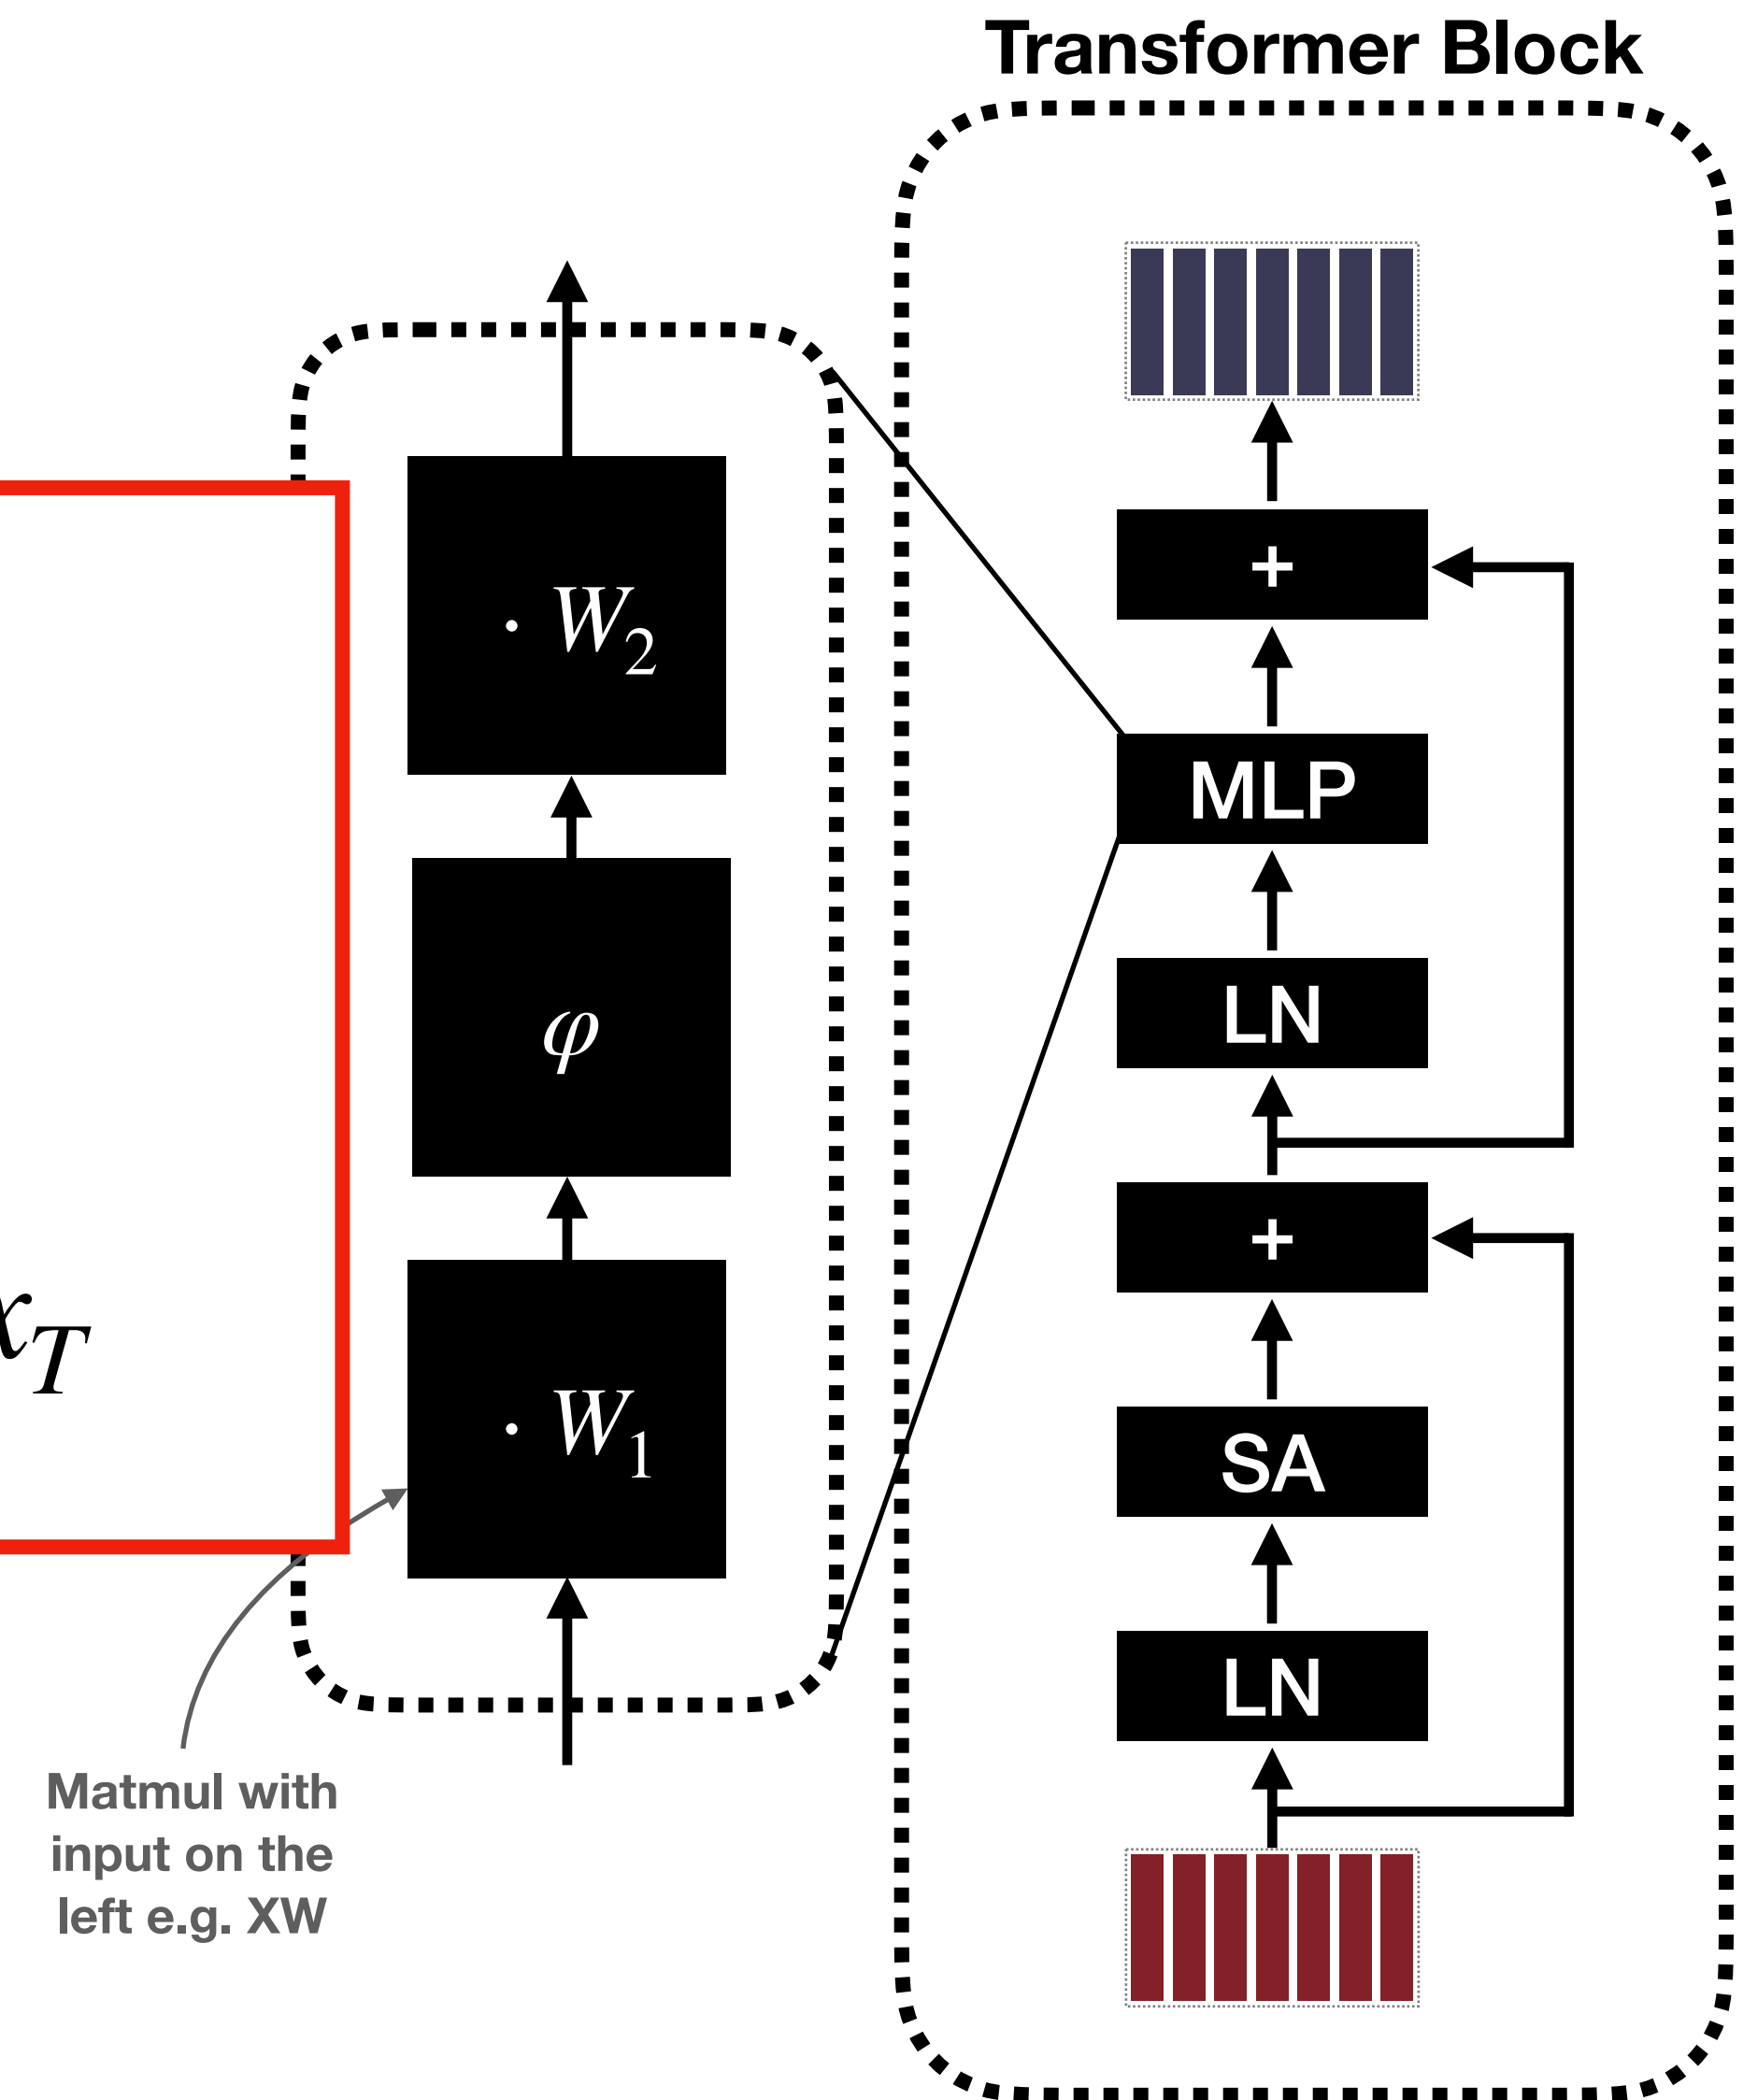
The same MLP is applied to each token:

$$MLP(X) = \varphi(XW_1)W_2$$



$$MLP(x_i) = \varphi(x_iW_1)W_2, \text{ for each token } x_1, \dots, x_T$$

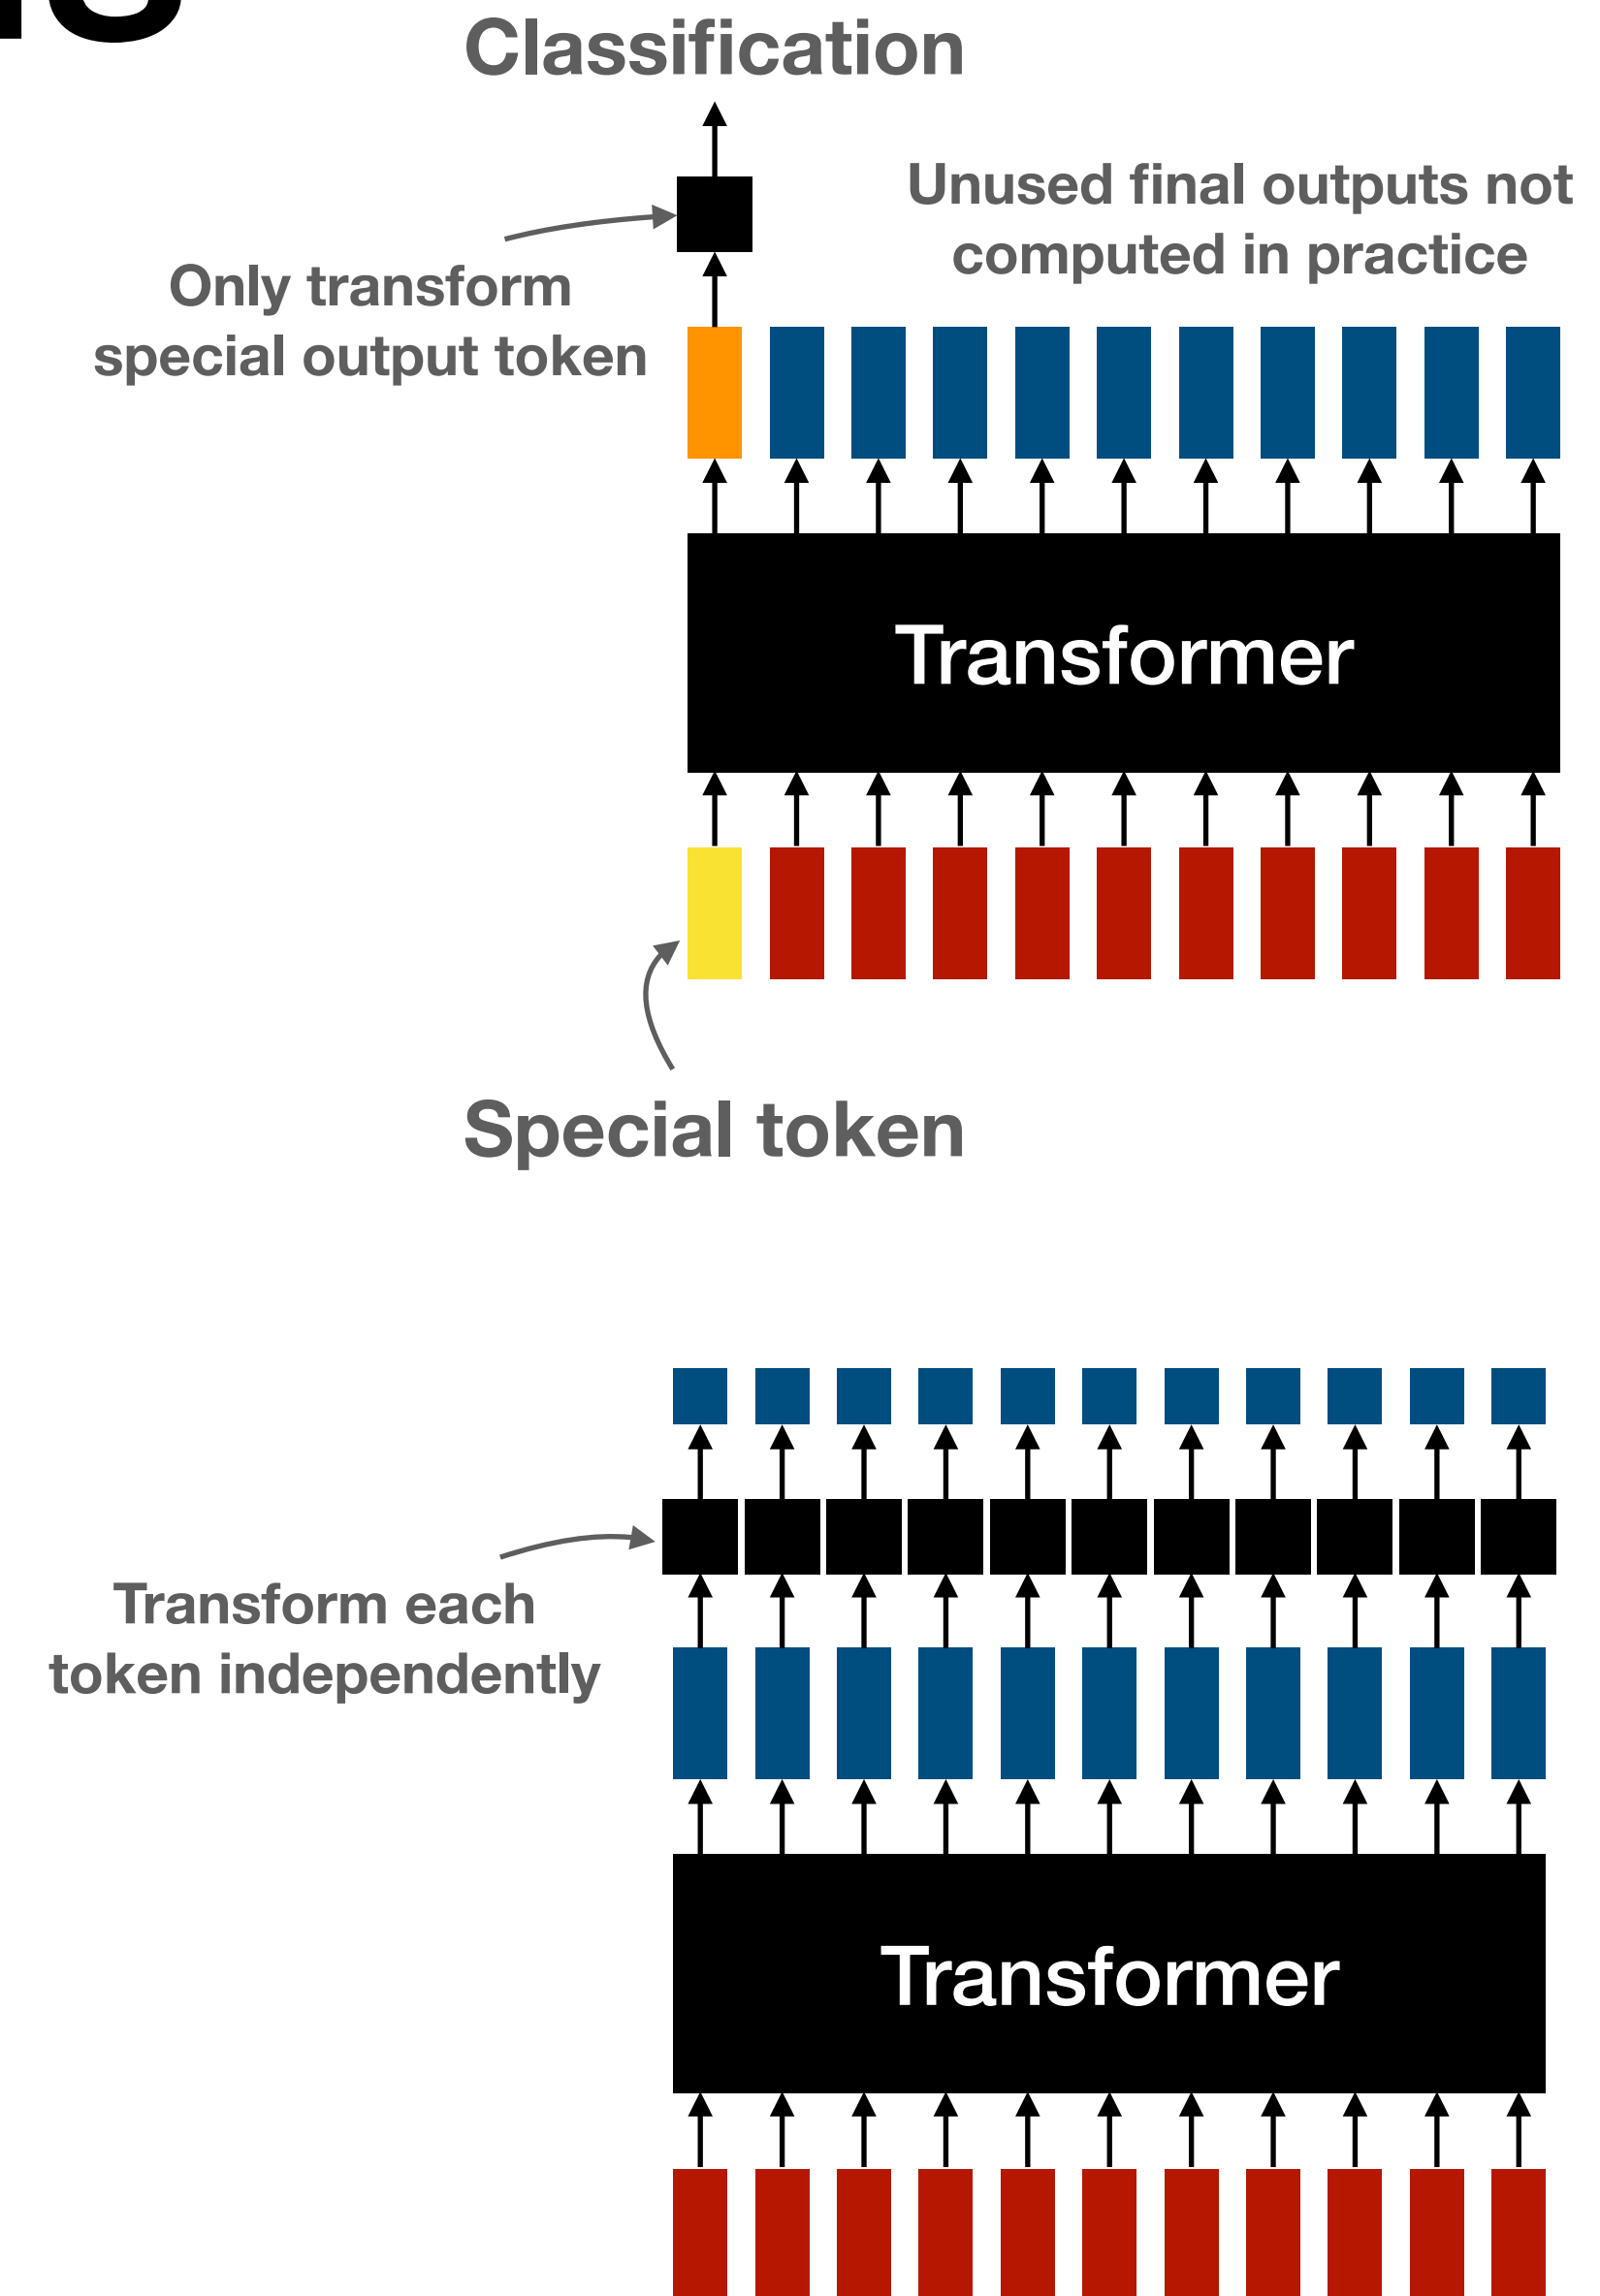
- Non-linearity  $\varphi$  in between (e.g., ReLU or GeLU)
- The model may also include learned bias terms



# Output Transformations

# Output Transformations

- We obtain the output from the final transformer block
- Output transformation is typically simple: linear transformation or a small MLP
- The specifics are highly dependent on the task:
  - **Single output** (e.g., sequence-level classification): apply an output transformation to a special task-specific input token or to the average of all tokens
  - **Multiple outputs** (e.g., per-token classification): apply an output transformation to each token independently

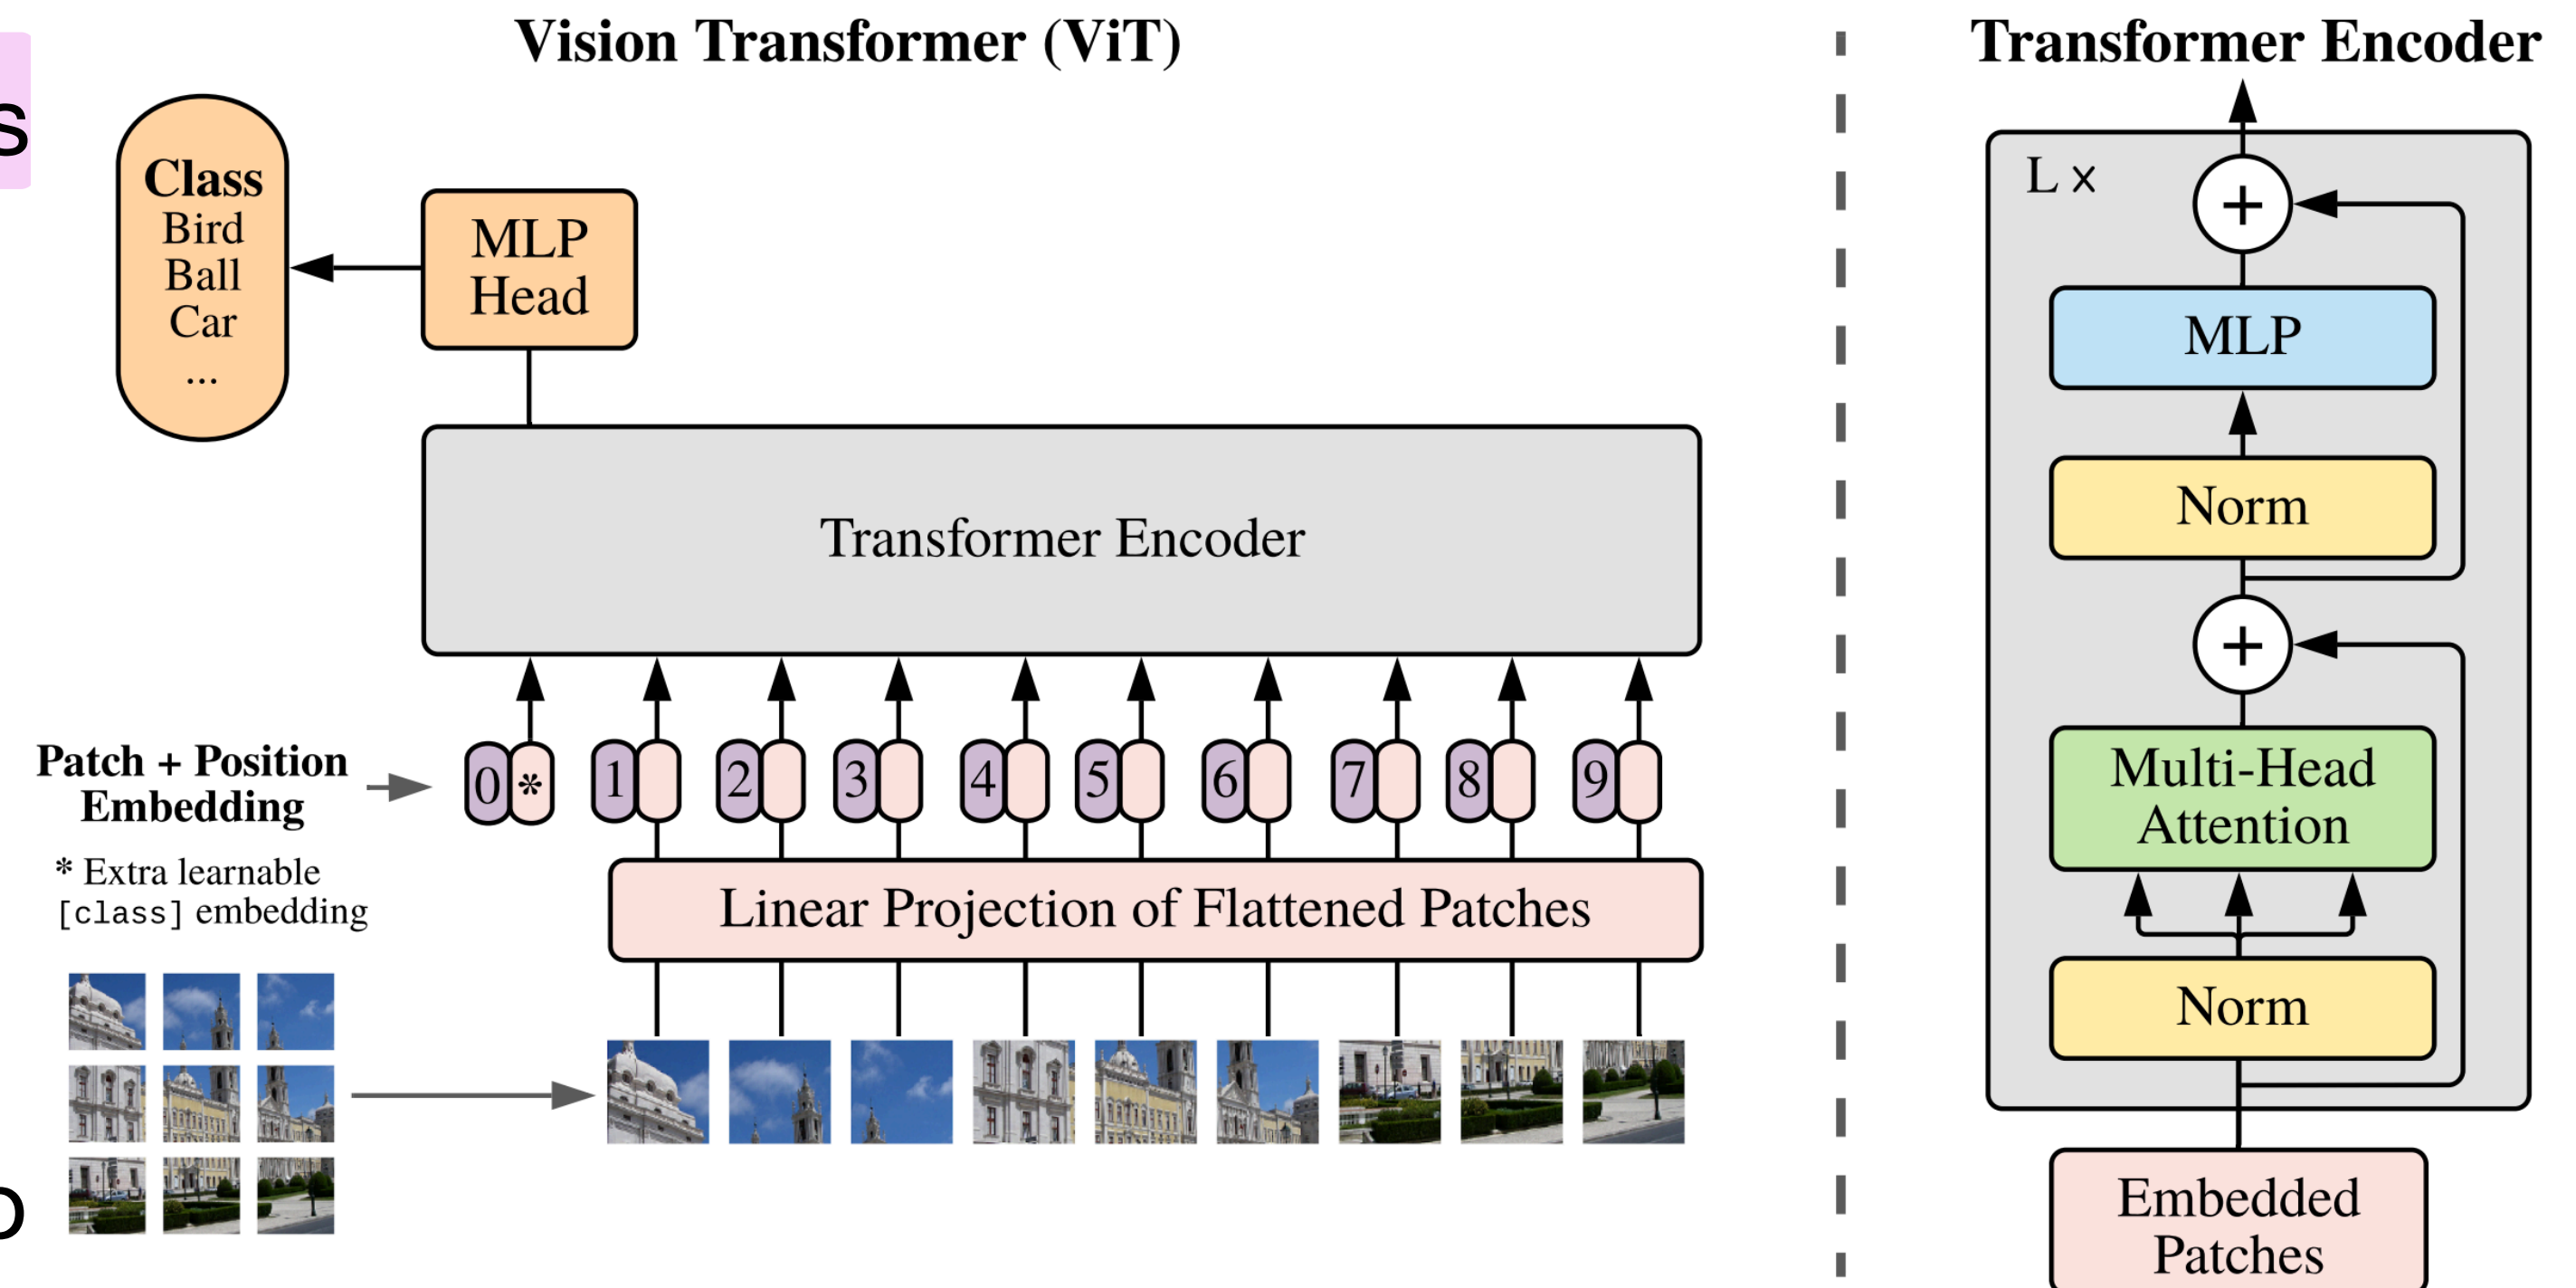


# Putting the pieces together: Vision Transformers



# Vision Transformer Architecture

- **Simple architecture:** number of features  $D$  is constant across all layers. There is no use of padding, pooling, or strides.
- Self-attention is **more general** than convolution and can express it
- The receptive field is the **whole image** after just one self-attention layer
- ViTs require more data than CNNs due to their reduced inductive bias in extracting local features
- However, ViTs become competitive with CNNs after **large-scale pretraining**



$$\mathbf{z}_0 = [\mathbf{x}_{\text{class}}; \mathbf{x}_p^1 \mathbf{E}; \mathbf{x}_p^2 \mathbf{E}; \dots; \mathbf{x}_p^N \mathbf{E}] + \mathbf{E}_{\text{pos}}, \quad \mathbf{E} \in \mathbb{R}^{(P^2 \cdot C) \times D}, \mathbf{E}_{\text{pos}} \in \mathbb{R}^{(N+1) \times D}$$


$$\mathbf{z}'_\ell = \text{MSA}(\text{LN}(\mathbf{z}_{\ell-1})) + \mathbf{z}_{\ell-1}, \quad \ell = 1 \dots L$$

$$\mathbf{z}_\ell = \text{MLP}(\text{LN}(\mathbf{z}'_\ell)) + \mathbf{z}'_\ell, \quad \ell = 1 \dots L$$

$$\mathbf{y} = \text{LN}(\mathbf{z}_L^0)$$

Source: An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale (ICLR 2020)

# What do ViTs learn: embedding layer



The first 28 principal components of the embedding layer applied on patches

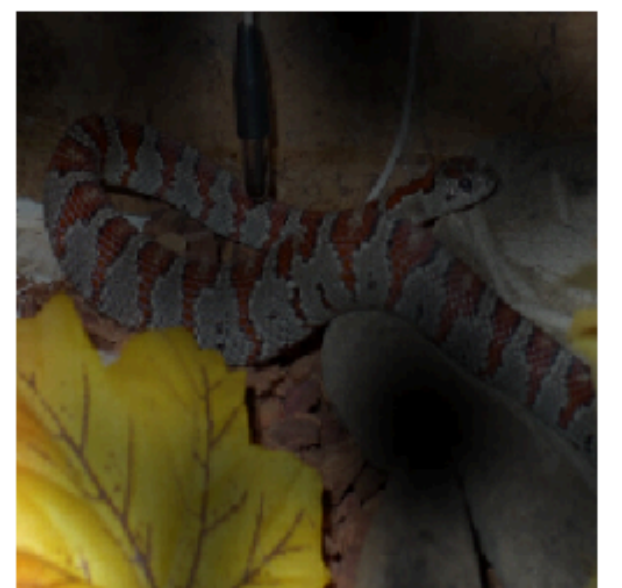
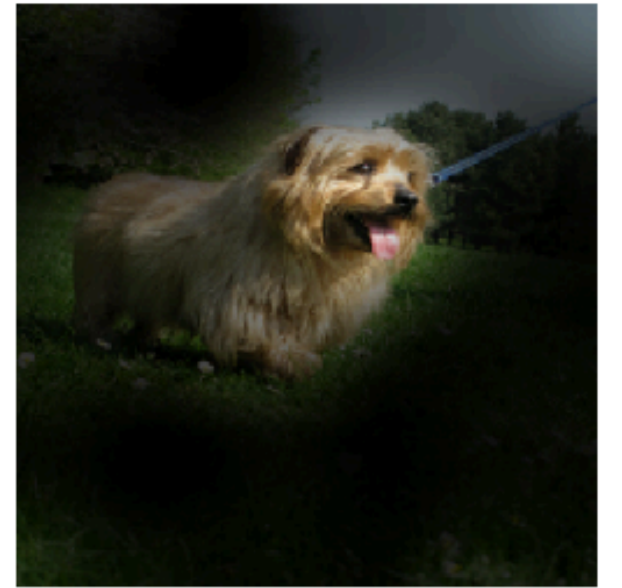
**Source:** An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale (ICLR 2020)

- The embedding layer: edge/color detectors similar to first-layer convolutions



# What do ViTs learn: attention


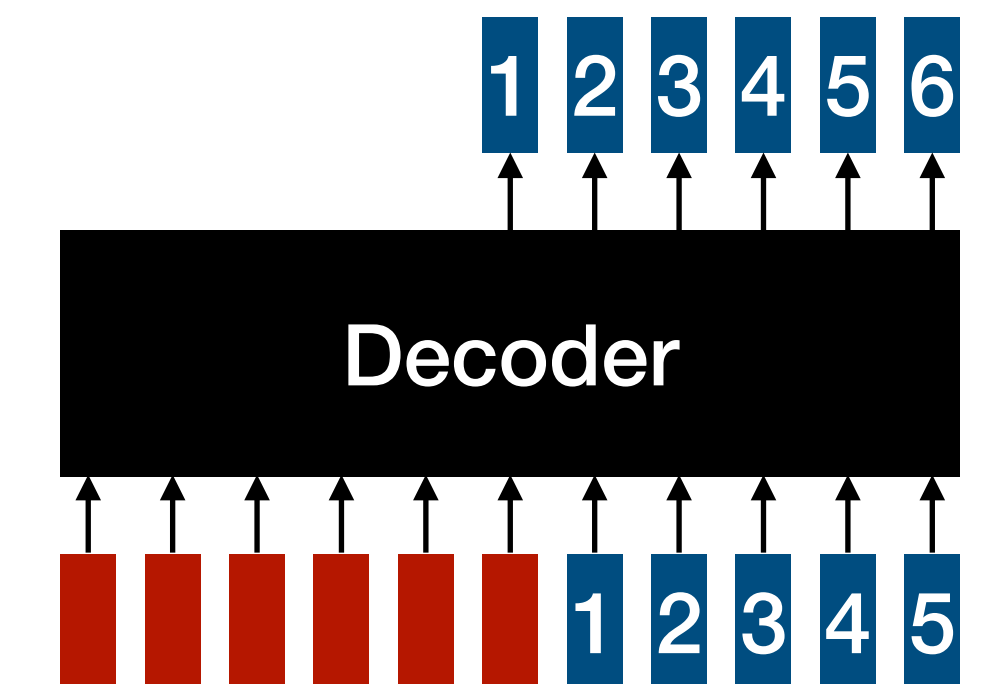
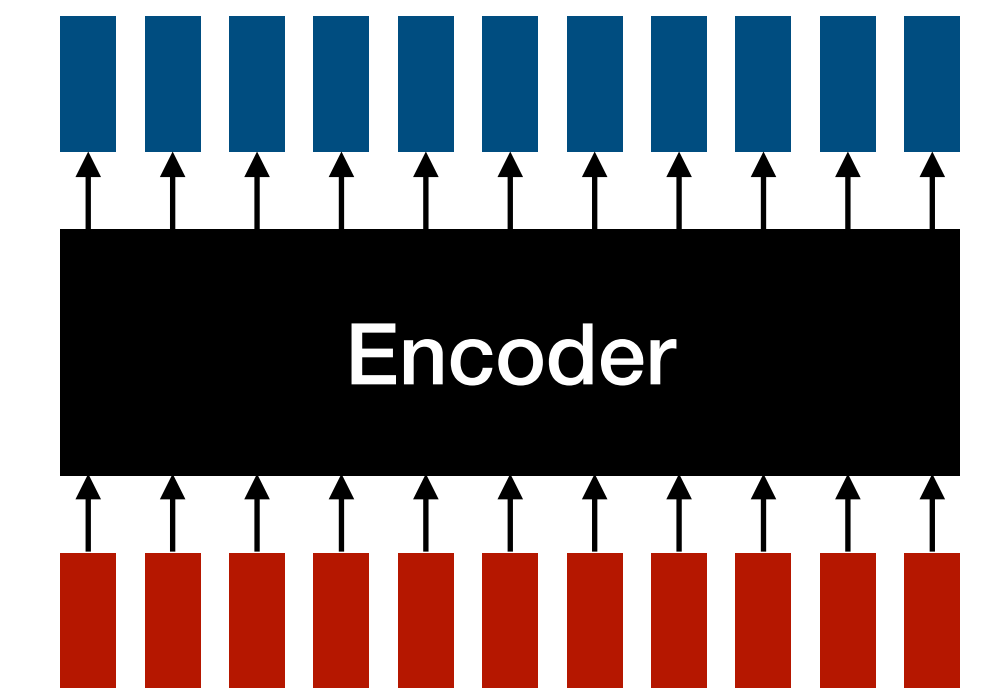
- The input-dependent attention weights can be visualized and manually inspected
- We show here one particular method known as Attention Rollout: where the attention weights are averaged across all heads and the resulting weight matrices of all layers are multiplied together
- This accounts for the mixing of attention across tokens through all layers
- In many cases, the model attends to image regions that are **semantically relevant** for classification



# The Big Picture and Takeaways

# The transformer architecture can be used in different ways

- **Encoders** (e.g., classification):
  - They produce a fixed output size and process all inputs simultaneously
- **Decoders** (e.g., ChatGPT):
  - **Auto-regressively sample** the next token as  $x_{t+1} \sim \text{softmax}(f(x_1, \dots, x_t))$  and use it as **new input token**
  - Capable of generating responses of arbitrary length
- **Encoder-decoder** (e.g., translation):
  - First encode the whole input (e.g., in one language) and then decode to token by token (e.g., in a different language)





# Transformers: Big Picture

- **Everything can be seen as a token, hence transformers are applicable across any modality**
- **CNNs can also be used for text processing, but transformers excel at capturing long-range dependencies** (as an example, the latest GPT-4 model can process up to 128k input tokens, equivalent to ~300 pages of text).
- **Self-attention scales quadratically with sequence length**, making it computationally expensive for large volumes of text or numerous patches—active area of research
- **However, self-attention is highly parallelizable**, which is advantageous for multi-GPU or multi-node training setups
- Transformers are now the **preferred method** for both text and vision applications
- **Emergent abilities at scale**: few-shot learning (aka in-context learning from a few example) and zero-shot learning (e.g., you can ask ChatGPT any question without prior training on the task)



# Recap

- **Transformers** iteratively map sequences to sequences using the self-attention mechanism
- The whole architecture is remarkably simple:
  - **Self-attention blocks** mix the information **between** tokens
  - **MLP blocks** mix the information **within** each token
- Transformers excel at modeling long-range dependencies
- Different architectures are possible (e.g., ChatGPT is decoder-only, but neural translation typically employs an encoder-decoder)
- Transformers have become a **universal architecture** for almost any type of data modality; they perform exceptionally well when given enough pretraining data

# Additional Resources

If you want to learn more about attention and transformers:

- **The Illustrated Transformer:** <https://jalammar.github.io/illustrated-transformer/> (a good step-by-step guide with detailed illustrations)
- **The blog of Lilian Weng (OpenAI):** <https://lilianweng.github.io/posts/2018-06-24-attention/> (from 2018 but covers well the history of the attention mechanism and its different versions)
- **CS231n: Deep Learning for Computer Vision (Stanford):** [http://cs231n.stanford.edu/slides/2023/lecture\\_9.pdf](http://cs231n.stanford.edu/slides/2023/lecture_9.pdf) (more on positional encodings, masked self-attention, general attention, discussion of recurrent neural networks)
- **Minimal implementation of GPT-2:** <https://github.com/karpathy/nanoGPT/> (some things are just clearer in code)